

## ✓ True vs Fake News Text Classification

### ✓ 3.1 Text Preprocessing, Tokenization, and Sequence Padding:

```
from google.colab import drive
drive.mount('/content/drive')
```

↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

#### ✓ Import Required Libraries:

```
# Import libraries for data manipulation
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Word cloud for data visualization
from wordcloud import WordCloud
```

```
# Regular expressions and string handling
import re
import string
```

```
# Natural Language Toolkit (NLTK) for text processing
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
```

```
# Download required NLTK resources
nltk.download('punkt')           # Tokenizer models
nltk.download('punkt_tab')
nltk.download('stopwords')       # Stopword list
nltk.download('wordnet')         # Lemmatizer dictionary
nltk.download('omw-1.4')        # Lemmatizer wordnet data
```

↗ [nltk\_data] Downloading package punkt to /root/nltk\_data...  
 [nltk\_data] Package punkt is already up-to-date!  
 [nltk\_data] Downloading package punkt\_tab to /root/nltk\_data...  
 [nltk\_data] Package punkt\_tab is already up-to-date!  
 [nltk\_data] Downloading package stopwords to /root/nltk\_data...  
 [nltk\_data] Package stopwords is already up-to-date!  
 [nltk\_data] Downloading package wordnet to /root/nltk\_data...  
 [nltk\_data] Package wordnet is already up-to-date!  
 [nltk\_data] Downloading package omw-1.4 to /root/nltk\_data...  
 [nltk\_data] Package omw-1.4 is already up-to-date!  
 True

```
!pip install -U numpy gensim
```

↗ Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (2.0.2)  
 Collecting numpy  
 Downloading numpy-2.2.5-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (62 kB)  
 \_\_\_\_\_ 60.6/60.6 kB 4.2 MB/s eta 0:00:00  
 Collecting gensim  
 Downloading gensim-4.3.3-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (8.1 kB)  
 Collecting numpy  
 Downloading numpy-1.26.4-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (61 kB)  
 \_\_\_\_\_ 61.0/61.0 kB 5.6 MB/s eta 0:00:00  
 Collecting scipy<1.14.0,>=1.7.0 (from gensim)  
 Downloading scipy-1.13.1-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (60 kB)  
 \_\_\_\_\_ 60.6/60.6 kB 5.8 MB/s eta 0:00:00  
 Requirement already satisfied: smart-open>=1.8.1 in /usr/local/lib/python3.11/dist-packages (from gensim) (7.1.0)  
 Requirement already satisfied: wrapt in /usr/local/lib/python3.11/dist-packages (from smart-open>=1.8.1->gensim) (1.17.2)  
 Downloading gensim-4.3.3-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (26.7 MB)  
 \_\_\_\_\_ 26.7/26.7 MB 84.1 MB/s eta 0:00:00  
 Downloading numpy-1.26.4-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (18.3 MB)  
 \_\_\_\_\_ 18.3/18.3 MB 95.5 MB/s eta 0:00:00  
 Downloading scipy-1.13.1-cp311-cp311-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (38.6 MB)

38.6/38.6 MB 18.5 MB/s eta 0:00:00

Installing collected packages: numpy, scipy, gensim

Attempting uninstall: numpy

Found existing installation: numpy 2.0.2

Uninstalling numpy-2.0.2:

Successfully uninstalled numpy-2.0.2

Attempting uninstall: scipy

Found existing installation: scipy 1.15.2

Uninstalling scipy-1.15.2:

Successfully uninstalled scipy-1.15.2

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the current problem. Upgrade pip to version 23.0 or newer to fix this. You may have to rebuild packages after upgrade. This problem is common in virtual environments created using pip.

tsfresh 0.21.0 requires scipy>=1.14.0; python\_version >= "3.10", but you have scipy 1.13.1 which is incompatible.

Successfully installed gensim-4.3.3 numpy-1.26.4 scipy-1.13.1

## Load the dataset:

# Load the dataset

file\_path = '/content/drive/MyDrive/AI ML/10.True vs. Fake News Dataset/truevsfakenews.csv'

data = pd.read\_csv(file\_path)

data.head()



	text	label
0	WASHINGTON (Reuters) - The Republican and Demo...	true
1	Women should get as far away from Oklahoma as ...	fake
2	Another huge crowd of Americans tuned in last ...	fake
3	Donald Trump is desperate to stop the investig...	fake
4	(Reuters) - Planned Parenthood, the U.S. medic...	true

## Clean the text:

# Drop missing values

data = data[['text', 'label']].dropna()

# Normalize labels to lowercase

data['label'] = data['label'].str.lower()

## Define Preprocessing Function

# Initialize lemmatizer and stopwords

lemmatizer = WordNetLemmatizer()

stop\_words = set(stopwords.words('english'))

# Contraction mapping

```
contractions_dict = {
    "don't": "do not", "doesn't": "does not", "didn't": "did not",
    "can't": "cannot", "won't": "will not", "shouldn't": "should not",
    "isn't": "is not", "aren't": "are not", "wasn't": "was not",
    "weren't": "were not", "hasn't": "has not", "haven't": "have not",
    "hadn't": "had not", "mightn't": "might not", "mustn't": "must not",
    "i'm": "i am", "you're": "you are", "he's": "he is", "she's": "she is",
    "it's": "it is", "we're": "we are", "they're": "they are",
    "i've": "i have", "you've": "you have", "we've": "we have",
    "they've": "they have", "i'll": "i will", "you'll": "you will",
    "he'll": "he will", "she'll": "she will", "we'll": "we will",
    "they'll": "they will"
}
```

# Compile regex pattern for contractions

```
contraction_pattern = re.compile(
    r'\b({})\b'.format('|'.join(re.escape(k) for k in contractions_dict.keys()))
    flags=re.IGNORECASE
)
```

# Function to fix Unicode apostrophe replacements like â€™

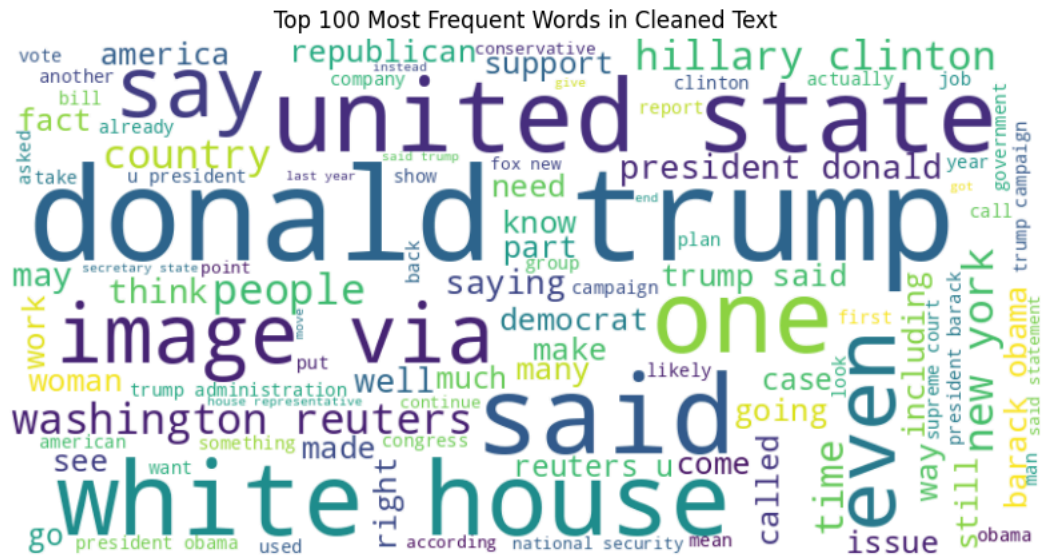
def fix\_unicode\_artifacts(text):

# Replace common Unicode misinterpretations



```
wordcloud = WordCloud(width=800, height=400, background_color='white', max_words=100).generate(all_words)
```

```
plt.figure(figsize=(10, 6))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Top 100 Most Frequent Words in Cleaned Text')
plt.show()
```



- ▼ Train/Test Split

```
from sklearn.model_selection import train_test_split
```

```
# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    data['cleaned_text'],
    data['label'],
    test_size=0.2,
    random_state=42,
    stratify=data['label']
)
```

- Tokenization

```
from tensorflow.keras.preprocessing.text import Tokenizer
import numpy as np
```

```
# Initialize tokenizer
tokenizer = Tokenizer(num_words=10000, oov_token='<OOV>')
tokenizer.fit_on_texts(X_train) # Only fit on training data
```

```
# Convert texts to sequences
X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)
```

- ✓ Padding Sequences Based on 95th Percentile

```
from tensorflow.keras.preprocessing.sequence import pad_sequences
```


```
# Calculate sequence lengths
seq_lengths = [len(seq) for seq in X_train_seq]
```

```
# Get 95th percentile of sequence lengths
max_len = int(np.percentile(seq_lengths, 95))
print(f"Padding sequences to maximum length: {max_len}")
```

```
# Pad sequences
```

```
X_train_pad = pad_sequences(X_train_seq, maxlen=max_len, padding='post', truncating='post')
X_test_pad = pad_sequences(X_test_seq, maxlen=max_len, padding='post', truncating='post')
```

```
# Check shapes
print(f"\nTraining data shape after padding: {X_train_pad.shape}")
print(f"Testing data shape after padding: {X_test_pad.shape}")
```

 Padding sequences to maximum length: 494


```
Training data shape after padding: (16000, 494)
Testing data shape after padding: (4000, 494)
```

## ✓ Label Encoding

```
from sklearn.preprocessing import LabelEncoder
```

```
# Encode string labels to integers
le = LabelEncoder()
y_train_enc = le.fit_transform(y_train)
y_test_enc = le.transform(y_test)

# Get number of classes (should be 2 for binary classification)
num_classes = len(le.classes_) # Output: 2
print(f"Number of classes: {num_classes}")
print(f"Class names: {list(le.classes_)}")
```

 Number of classes: 2  
Class names: ['fake', 'true']

## ✓ 3.2 Model Building and Training

### ✓ Import Required Libraries

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, LSTM, Dense
from tensorflow.keras.initializers import Constant
import gensim.downloader as api
import numpy as np
```

### ✓ MODEL 1: Simple RNN with Trainable Embedding

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense
```

```
model_rnn = Sequential()
```

```
# Embedding Layer
model_rnn.add(Embedding(
    input_dim=10000,
    output_dim=128,
    input_length=max_len
))
```


```
# Recurrent Layer
model_rnn.add(SimpleRNN(units=64))
```

```
# Output Layer
model_rnn.add(Dense(1, activation='sigmoid'))
```

```
# Compile Model
model_rnn.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy']
)
```

```
# Build the model manually by specifying the input shape
model_rnn.build(input_shape=(None, X_train_pad.shape[1]))
```

```
# Print Model Summary
model_rnn.summary()
```

 /usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embedding.py:90: UserWarning: Argument `input\_length` is deprecated. Just warnings.warn(  
Model: "sequential"

Layer (type)	Output Shape	Param #
embedding ( <a href="#">Embedding</a> )	(None, 494, 128)	1,280,000
simple_rnn ( <a href="#">SimpleRNN</a> )	(None, 64)	12,352
dense ( <a href="#">Dense</a> )	(None, 1)	65

Total params: 1,292,417 (4.93 MB)  
Trainable params: 1,292,417 (4.93 MB)

## ✓ MODEL 2: LSTM with Trainable Embedding

```
from tensorflow.keras.layers import LSTM

model_lstm = Sequential()

# Embedding Layer
model_lstm.add(Embedding(
    input_dim=10000,
    output_dim=128,
    input_length=max_len
))


# LSTM Layer
model_lstm.add(LSTM(units=64))

# Output Layer
model_lstm.add(Dense(1, activation='sigmoid'))

# Compile
model_lstm.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy']
)

# Build the model
model_lstm.build(input_shape=(None, X_train_pad.shape[1]))

model_lstm.summary()
```

 Model: "sequential\_1"


Layer (type)	Output Shape	Param #
embedding_1 ( <a href="#">Embedding</a> )	(None, 494, 128)	1,280,000
lstm ( <a href="#">LSTM</a> )	(None, 64)	49,408
dense_1 ( <a href="#">Dense</a> )	(None, 1)	65

Total params: 1,329,473 (5.07 MB)  
Trainable params: 1,329,473 (5.07 MB)  
Non-trainable params: 0 (0.00 B)

## ✓ MODEL 3: LSTM with Pretrained Word2Vec Embeddings

### ✓ Download and load a pretrained Word2Vec model

```
embedding_model = api.load("glove-wiki-gigaword-50") # 50-dimensional vectors
embedding_dim = 50
```

 [=====] 100.0% 66.0/66.0MB downloaded

### ✓ Create Embedding Matrix

```

vocab_size = min(len(tokenizer.word_index) + 1, 11000) # Cap at 10,000 words
embedding_matrix = np.zeros((vocab_size, embedding_dim))

for word, i in tokenizer.word_index.items():
    if i >= 10000:
        continue
    if word in embedding_model:
        embedding_vector = embedding_model[word]
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector

```

### Define LSTM Model Using Word2Vec Embeddings

```

from tensorflow.keras.layers import Embedding
from tensorflow.keras.initializers import Constant

model_lstm_word2vec = Sequential()

# Embedding Layer with Pretrained Weights
model_lstm_word2vec.add(Embedding(
    input_dim=vocab_size,
    output_dim=embedding_dim,
    input_length=max_len,
    embeddings_initializer=Constant(embedding_matrix), # Use pretrained weights
    trainable=False # Freeze embeddings unless fine-tuning
))

# LSTM Layer
model_lstm_word2vec.add(LSTM(units=64))

# Output Layer
model_lstm_word2vec.add(Dense(1, activation='sigmoid'))

# Compile
model_lstm_word2vec.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy']
)

# Build the model manually to initialize layers
model_lstm_word2vec.build(input_shape=(None, X_train_pad.shape[1]))

model_lstm_word2vec.summary()

```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
embedding_6 ( <a href="#">Embedding</a> )	(None, 494, 50)	550,000
lstm_5 ( <a href="#">LSTM</a> )	(None, 64)	29,440
dense_6 ( <a href="#">Dense</a> )	(None, 1)	65

Total params: 579,505 (2.21 MB)  
 Trainable params: 29,505 (115.25 KB)  
 Non-trainable params: 550,000 (2.10 MB)

### 3.3 Model Training and Evaluation:

#### Define Callbacks

```

from tensorflow.keras.callbacks import EarlyStopping

early_stop = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True
)

```

## ✓ Train Simple RNN

```
history_rnn = model_rnn.fit(
    X_train_pad,
    y_train_enc,
    validation_split=0.2,
    epochs=20,
    batch_size=32,
    callbacks=[early_stop],
    verbose=1
)
```

```
Epoch 1/20
400/400 ————— 15s 37ms/step - accuracy: 0.5266 - loss: 0.6704 - val_accuracy: 0.4956 - val_loss: 0.6858
Epoch 2/20
400/400 ————— 15s 37ms/step - accuracy: 0.5368 - loss: 0.6604 - val_accuracy: 0.4959 - val_loss: 0.6859
Epoch 3/20
400/400 ————— 15s 38ms/step - accuracy: 0.5373 - loss: 0.6603 - val_accuracy: 0.5344 - val_loss: 0.6848
Epoch 4/20
400/400 ————— 21s 38ms/step - accuracy: 0.5311 - loss: 0.6590 - val_accuracy: 0.5312 - val_loss: 0.6909
Epoch 5/20
400/400 ————— 20s 38ms/step - accuracy: 0.5356 - loss: 0.6598 - val_accuracy: 0.5316 - val_loss: 0.6865
Epoch 6/20
400/400 ————— 15s 38ms/step - accuracy: 0.5370 - loss: 0.6603 - val_accuracy: 0.4906 - val_loss: 0.6884
Epoch 7/20
400/400 ————— 20s 37ms/step - accuracy: 0.5371 - loss: 0.6605 - val_accuracy: 0.5319 - val_loss: 0.6877
Epoch 8/20
400/400 ————— 21s 38ms/step - accuracy: 0.5405 - loss: 0.6593 - val_accuracy: 0.5312 - val_loss: 0.6877
```

## ✓ Train LSTM

```
history_lstm = model_lstm.fit(
    X_train_pad,
    y_train_enc,
    validation_split=0.2,
    epochs=20,
    batch_size=32,
    callbacks=[early_stop],
    verbose=1
)
```

```
Epoch 1/20
400/400 ————— 11s 20ms/step - accuracy: 0.5334 - loss: 0.6899 - val_accuracy: 0.5456 - val_loss: 0.6722
Epoch 2/20
400/400 ————— 7s 18ms/step - accuracy: 0.5387 - loss: 0.6716 - val_accuracy: 0.5000 - val_loss: 0.6690
Epoch 3/20
400/400 ————— 8s 21ms/step - accuracy: 0.5617 - loss: 0.6618 - val_accuracy: 0.7881 - val_loss: 0.4676
Epoch 4/20
400/400 ————— 10s 20ms/step - accuracy: 0.8980 - loss: 0.3155 - val_accuracy: 0.9684 - val_loss: 0.1089
Epoch 5/20
400/400 ————— 10s 18ms/step - accuracy: 0.9870 - loss: 0.0569 - val_accuracy: 0.9669 - val_loss: 0.1048
Epoch 6/20
400/400 ————— 8s 19ms/step - accuracy: 0.9923 - loss: 0.0377 - val_accuracy: 0.9787 - val_loss: 0.0870
Epoch 7/20
400/400 ————— 8s 21ms/step - accuracy: 0.9956 - loss: 0.0211 - val_accuracy: 0.9816 - val_loss: 0.0630
Epoch 8/20
400/400 ————— 9s 18ms/step - accuracy: 0.9985 - loss: 0.0067 - val_accuracy: 0.9825 - val_loss: 0.0655
Epoch 9/20
400/400 ————— 10s 18ms/step - accuracy: 0.9995 - loss: 0.0034 - val_accuracy: 0.9841 - val_loss: 0.0692
Epoch 10/20
400/400 ————— 8s 21ms/step - accuracy: 0.9998 - loss: 0.0018 - val_accuracy: 0.9850 - val_loss: 0.0589
Epoch 11/20
400/400 ————— 10s 20ms/step - accuracy: 0.9995 - loss: 0.0034 - val_accuracy: 0.9881 - val_loss: 0.0618
Epoch 12/20
400/400 ————— 9s 18ms/step - accuracy: 0.9997 - loss: 0.0020 - val_accuracy: 0.9866 - val_loss: 0.0731
Epoch 13/20
400/400 ————— 10s 18ms/step - accuracy: 1.0000 - loss: 3.1964e-04 - val_accuracy: 0.9819 - val_loss: 0.1102
Epoch 14/20
400/400 ————— 8s 21ms/step - accuracy: 0.9980 - loss: 0.0073 - val_accuracy: 0.9787 - val_loss: 0.0948
Epoch 15/20
400/400 ————— 10s 20ms/step - accuracy: 0.9990 - loss: 0.0028 - val_accuracy: 0.9831 - val_loss: 0.0787
```

## ✓ Train LSTM with Word2Vec

```
history_word2vec = model_lstm_word2vec.fit(
    X_train_pad,
```



```

x_train_enc,
y_train_enc,
validation_split=0.2,
epochs=20,
batch_size=32,
callbacks=[early_stop],
verbose=1
)

```

```

Epoch 1/20
400/400 — 9s 18ms/step - accuracy: 0.5689 - loss: 0.6584 - val_accuracy: 0.6194 - val_loss: 0.5873
Epoch 2/20
400/400 — 8s 19ms/step - accuracy: 0.7218 - loss: 0.4937 - val_accuracy: 0.9328 - val_loss: 0.2881
Epoch 3/20
400/400 — 7s 18ms/step - accuracy: 0.7648 - loss: 0.4771 - val_accuracy: 0.5534 - val_loss: 0.6689
Epoch 4/20
400/400 — 10s 18ms/step - accuracy: 0.6393 - loss: 0.5924 - val_accuracy: 0.9137 - val_loss: 0.2593
Epoch 5/20
400/400 — 12s 22ms/step - accuracy: 0.9129 - loss: 0.2562 - val_accuracy: 0.9331 - val_loss: 0.2244
Epoch 6/20
400/400 — 7s 18ms/step - accuracy: 0.9375 - loss: 0.2140 - val_accuracy: 0.9344 - val_loss: 0.2367
Epoch 7/20
400/400 — 10s 17ms/step - accuracy: 0.9227 - loss: 0.2494 - val_accuracy: 0.9247 - val_loss: 0.2353
Epoch 8/20
400/400 — 8s 20ms/step - accuracy: 0.9223 - loss: 0.2388 - val_accuracy: 0.9337 - val_loss: 0.2214
Epoch 9/20
400/400 — 7s 17ms/step - accuracy: 0.9428 - loss: 0.2078 - val_accuracy: 0.9481 - val_loss: 0.1996
Epoch 10/20
400/400 — 11s 18ms/step - accuracy: 0.9485 - loss: 0.1939 - val_accuracy: 0.9319 - val_loss: 0.2269
Epoch 11/20
400/400 — 10s 17ms/step - accuracy: 0.9393 - loss: 0.2082 - val_accuracy: 0.9294 - val_loss: 0.2268
Epoch 12/20
400/400 — 11s 19ms/step - accuracy: 0.9353 - loss: 0.2155 - val_accuracy: 0.9500 - val_loss: 0.1950
Epoch 13/20
400/400 — 10s 19ms/step - accuracy: 0.9247 - loss: 0.2452 - val_accuracy: 0.9475 - val_loss: 0.1927
Epoch 14/20
400/400 — 10s 19ms/step - accuracy: 0.9526 - loss: 0.1810 - val_accuracy: 0.9569 - val_loss: 0.1791
Epoch 15/20
400/400 — 10s 18ms/step - accuracy: 0.9543 - loss: 0.1750 - val_accuracy: 0.9575 - val_loss: 0.1715
Epoch 16/20
400/400 — 8s 19ms/step - accuracy: 0.9549 - loss: 0.1734 - val_accuracy: 0.9569 - val_loss: 0.1669
Epoch 17/20
400/400 — 11s 21ms/step - accuracy: 0.8995 - loss: 0.3141 - val_accuracy: 0.6031 - val_loss: 0.6326
Epoch 18/20
400/400 — 9s 19ms/step - accuracy: 0.7323 - loss: 0.5014 - val_accuracy: 0.9228 - val_loss: 0.2373
Epoch 19/20
400/400 — 7s 18ms/step - accuracy: 0.8281 - loss: 0.3568 - val_accuracy: 0.9397 - val_loss: 0.2018
Epoch 20/20
400/400 — 8s 19ms/step - accuracy: 0.9401 - loss: 0.1987 - val_accuracy: 0.9463 - val_loss: 0.1891

```

## ▼ Plot Training vs Validation Accuracy & Loss

```

import matplotlib.pyplot as plt

def plot_history(history, model_name):
    plt.figure(figsize=(12, 4))

    # Accuracy
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title(f'{model_name} - Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend()


    # Loss
    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'], label='Train Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
    plt.title(f'{model_name} - Loss')
    plt.ylabel('Loss')
    plt.xlabel('Epoch')
    plt.legend()

    plt.tight_layout()
    plt.show()


```

## ✓ Plot Each Model's History


```
plot_history(history_rnn, "Simple RNN")
```



```
plot_history(history_lstm, "LSTM")
```



```
plot_history(history_word2vec, "LSTM + Word2Vec")
```



## ✓ Evaluate Models on Test Set

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import numpy as np
```

```
def evaluate_model(model, X_test, y_test, model_name):
    y_pred = model.predict(X_test)

    # Binary classification: threshold at 0.5
    y_pred_classes = (y_pred > 0.5).astype(int).flatten()

    acc = accuracy_score(y_test, y_pred_classes)
    print(f"\n{model_name} - Test Accuracy: {acc:.4f}")

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred_classes)
    print(f"\nConfusion Matrix:\n{cm}")

    # Classification Report
    cr = classification_report(y_test, y_pred_classes, target_names=le.classes_)
    print(f"\nClassification Report:\n{cr}")
```

## ✓ Run Evaluation for All Models

```
evaluate_model(model_rnn, X_test_pad, y_test_enc, "Simple RNN")
```

↗ 125/125 ————— 4s 21ms/step

Simple RNN - Test Accuracy: 0.5333

Confusion Matrix:

```
[[1972  28]
 [1839 161]]
```

Classification Report:

	precision	recall	f1-score	support
fake	0.52	0.99	0.68	2000
true	0.85	0.08	0.15	2000
accuracy			0.53	4000
macro avg	0.68	0.53	0.41	4000
weighted avg	0.68	0.53	0.41	4000

```
evaluate_model(model_lstm, X_test_pad, y_test_enc, "LSTM")
```

↗ 125/125 ————— 2s 11ms/step

LSTM - Test Accuracy: 0.9815

Confusion Matrix:

```
[[1957  43]
 [ 31 1969]]
```

Classification Report:

	precision	recall	f1-score	support
fake	0.98	0.98	0.98	2000
true	0.98	0.98	0.98	2000
accuracy			0.98	4000
macro avg	0.98	0.98	0.98	4000
weighted avg	0.98	0.98	0.98	4000

```
evaluate_model(model_lstm_word2vec, X_test_pad, y_test_enc, "LSTM + Word2Vec")
```

↗ 125/125 ————— 1s 7ms/step

LSTM + Word2Vec - Test Accuracy: 0.9483

Confusion Matrix:

```
[[1821 179]
 [ 28 1972]]
```

```

Classification Report:
              precision    recall  f1-score   support

    fake       0.98        0.91        0.95        2000
    true       0.92        0.99        0.95        2000

 accuracy          0.95          0.95          0.95        4000
 macro avg       0.95        0.95        0.95        4000
 weighted avg    0.95        0.95        0.95        4000

```

## ▼ Compare Model Performances

```

from prettytable import PrettyTable
from sklearn.metrics import accuracy_score

table = PrettyTable()
table.field_names = ["Model", "Test Accuracy"]

for model, name in zip(
    [model_rnn, model_lstm, model_lstm_word2vec],
    ["Simple RNN", "LSTM", "LSTM + Word2Vec"]
):
    y_pred = model.predict(X_test_pad)

    # For binary classification using sigmoid output
    y_pred_classes = (y_pred > 0.5).astype(int).flatten()

    acc = accuracy_score(y_test_enc, y_pred_classes)
    table.add_row([name, f"{acc:.4f}"])

print("\nModel Comparison:")
print(table)

```

```

125/125 ————— 1s 11ms/step
125/125 ————— 1s 9ms/step
125/125 ————— 1s 9ms/step

```

```

Model Comparison:
+-----+-----+
| Model | Test Accuracy |
+-----+-----+
| Simple RNN | 0.5333 |
| LSTM | 0.9815 |
| LSTM + Word2Vec | 0.9483 |
+-----+-----+

```

## ▼ Save Models

```

# Save Simple RNN Model
model_rnn.save("/content/drive/MyDrive/AI ML/simple_rnn_model.h5")

```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. Please use `model.save(format='tf')` or `keras.saving.save\_model(model, format='tf')` instead.

```

# Save LSTM Model
model_lstm.save("/content/drive/MyDrive/AI ML/lstm_model.h5")

```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. Please use `model.save(format='tf')` or `keras.saving.save\_model(model, format='tf')` instead.

```

# Save LSTM + Word2Vec Model
model_lstm_word2vec.save("/content/drive/MyDrive/AI ML/lstm_word2vec_model.h5")

```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. Please use `model.save(format='tf')` or `keras.saving.save\_model(model, format='tf')` instead.

## ▼ Load Required Libraries again

```
!pip install gradio
```

```
Collecting gradio
  Downloading gradio-5.29.0-py3-none-any.whl.metadata (16 kB)
Collecting aiofiles<25.0,>=22.0 (from gradio)
  Downloading aiofiles-24.1.0-py3-none-any.whl.metadata (10 kB)
Requirement already satisfied: anyio<5.0,>=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.9.0)
Collecting fastapi<1.0,>=0.115.2 (from gradio)
  Downloading fastapi-0.115.12-py3-none-any.whl.metadata (27 kB)
Collecting ffmpy (from gradio)
  Downloading ffmpy-0.5.0-py3-none-any.whl.metadata (3.0 kB)
Collecting gradio-client==1.10.0 (from gradio)
  Downloading gradio_client-1.10.0-py3-none-any.whl.metadata (7.1 kB)
Collecting groovy~=0.1 (from gradio)
  Downloading groovy-0.1.2-py3-none-any.whl.metadata (6.1 kB)
Requirement already satisfied: httpx>=0.24.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.28.1)
Requirement already satisfied: huggingface-hub>=0.28.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.31.1)
Requirement already satisfied: jinja2<4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.1.6)
Requirement already satisfied: markupsafe<4.0,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.0.2)
Requirement already satisfied: numpy<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (1.26.4)
Requirement already satisfied: orjson~=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.10.18)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from gradio) (24.2)
Requirement already satisfied: pandas<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.2.2)
Requirement already satisfied: pillow<12.0,>=8.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (11.2.1)
Requirement already satisfied: pydantic<2.12,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.11.4)
Collecting pydub (from gradio)
  Downloading pydub-0.25.1-py3-none-any.whl.metadata (1.4 kB)
Collecting python-multipart>=0.0.18 (from gradio)
  Downloading python_multipart-0.0.20-py3-none-any.whl.metadata (1.8 kB)
Requirement already satisfied: pyyaml<7.0,>=5.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (6.0.2)
Collecting ruff>=0.9.3 (from gradio)
  Downloading ruff-0.11.9-py3-none-manylinux_2_17_x86_64_musl.manylinux2014_x86_64.whl.metadata (25 kB)
Collecting safehttpx<0.2.0,>=0.1.6 (from gradio)
  Downloading safehttpx-0.1.6-py3-none-any.whl.metadata (4.2 kB)
Collecting semantic-version~=2.0 (from gradio)
  Downloading semantic_version-2.10.0-py2.py3-none-any.whl.metadata (9.7 kB)
Collecting starlette<1.0,>=0.40.0 (from gradio)
  Downloading starlette-0.46.2-py3-none-any.whl.metadata (6.2 kB)
Collecting tomlkit<0.14.0,>=0.12.0 (from gradio)
  Downloading tomlkit-0.13.2-py3-none-any.whl.metadata (2.7 kB)
Requirement already satisfied: typer<1.0,>=0.12 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.15.3)
Requirement already satisfied: typing-extensions~=4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.13.2)
Collecting uvicorn>=0.14.0 (from gradio)
  Downloading uvicorn-0.34.2-py3-none-any.whl.metadata (6.5 kB)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.10.0->gradio) (2025.3.2)
Requirement already satisfied: websockets<16.0,>=10.0 in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.10.0->gradio) (13.1)
Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.11/dist-packages (from anyio<5.0,>=3.0->gradio) (3.10)
Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.11/dist-packages (from anyio<5.0,>=3.0->gradio) (1.3.1)
Requirement already satisfied: certifi in /usr/local/lib/python3.11/dist-packages (from httpx>=0.24.1->gradio) (2025.4.26)
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.11/dist-packages (from httpx>=0.24.1->gradio) (1.0.9)
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.11/dist-packages (from httpcore==1.*->httpx>=0.24.1->gradio) (0.16)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (3.18.0)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (2.32.3)
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (4.67.1)
Requirement already satisfied: hf-xet<2.0.0,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (1.1.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2.9)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.2)
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio) (0.7.0)
```

```
import gradio as gr
import numpy as np
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from tensorflow.keras.models import load_model
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
# Ensure NLTK resources are available
nltk.download('stopwords')
nltk.download('wordnet')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
True
```

## ✓ Redefine Preprocessing Function

```
# Initialize lemmatizer and stopwords
lemmatizer = WordNetLemmatizer()
stop_words = set(stopwords.words('english'))

# Improved contraction mapping
contractions_dict = {
    "don't": "do not", "doesn't": "does not", "didn't": "did not",
    "can't": "cannot", "won't": "will not", "shouldn't": "should not",
    "isn't": "is not", "aren't": "are not", "wasn't": "was not",
    "weren't": "were not", "hasn't": "has not", "haven't": "have not",
    "hadn't": "had not", "mightn't": "might not", "mustn't": "must not",
    "i'm": "i am", "you're": "you are", "he's": "he is", "she's": "she is",
    "it's": "it is", "we're": "we are", "they're": "they are",
    "i've": "i have", "you've": "you have", "we've": "we have",
    "they've": "they have", "i'll": "i will", "you'll": "you will",
    "he'll": "he will", "she'll": "she will", "we'll": "we will",
    "they'll": "they will"
}

# Compile regex pattern for contractions
contraction_pattern = re.compile(
    r'\b({})\b'.format('|'.join(re.escape(k) for k in contractions_dict.keys())),
    flags=re.IGNORECASE
)

# Fix Unicode artifacts
def fix_unicode_artifacts(text):
    return text.replace("â€œ", "").replace("â€œ", "").replace("â€œ", "").replace("â€œ", "-").replace("Â", "")

# Expand contractions
def expand_contractions(text):
    def replace(match):
        word = match.group(0).lower()
        return contractions_dict.get(word, word)
    return contraction_pattern.sub(replace, text)

# Full preprocessing function
def preprocess_input(text):
    text = fix_unicode_artifacts(text)
    text = text.lower()
    text = expand_contractions(text)
    text = re.sub(r'https?://\S+|www.\S+', '', text)
    text = re.sub(r'@\w+|#', '', text)
    text = re.sub(r'^\w\s\']', '', text)
    text = re.sub(r'\d+', '', text)
    tokens = text.split()
    tokens = [lemmatizer.lemmatize(word) for word in tokens if word not in stop_words]
    return ' '.join(tokens)
```

## ✓ Load Saved Models

```
from tensorflow.keras.models import load_model

# Ask user for the directory where models are saved
save_path = input("Enter the directory path where your models are saved (e.g., /content/drive/MyDrive/AI ML): ")

# Ensure the path doesn't end with a slash for consistency
save_path = save_path.rstrip('/')

# Load models from user-specified path
loaded_rnn = load_model(f"{save_path}/simple_rnn_model.h5")
loaded_lstm = load_model(f"{save_path}/lstm_model.h5")
loaded_lstm_word2vec = load_model(f"{save_path}/lstm_word2vec_model.h5")

print("\n Models loaded successfully!")
```

Enter the directory path where your models are saved (e.g., /content/drive/MyDrive/AI ML): /content/drive/MyDrive/AI ML  
 WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you t  
 WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you t  
 WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you t

Models loaded successfully!

## ▼ Define Prediction Function

```
# Global tokenizer and max_len used during training
global_tokenizer = tokenizer # Use the one from earlier
global_max_len = X_train_pad.shape[1]

# Label encoder
le_classes = le.classes_ # e.g., ['fake', 'true']

def predict_news(news_text, model_choice):
    # Preprocess input
    cleaned_text = preprocess_input(news_text)

    # Tokenize and pad
    sequence = global_tokenizer.texts_to_sequences([cleaned_text])
    padded_seq = pad_sequences(sequence, maxlen=global_max_len, padding='post', truncating='post')

    # Choose model
    if model_choice == "Simple RNN":
        model = loaded_rnn
    elif model_choice == "LSTM":
        model = loaded_lstm
    else:
        model = loaded_lstm_word2vec

    # Predict
    prediction = model.predict(padded_seq, verbose=0)
    predicted_label = le_classes[np.argmax(prediction)]

    # Return styled HTML output
    if predicted_label == 'true':
        return "<div style='padding:20px; background-color:#d4edda; color:#155724; font-size:20px; border-radius:8px; text-align:center;'>✔️"
    else:
        return "<div style='padding:20px; background-color:#f8d7da; color:#721c24; font-size:20px; border-radius:8px; text-align:center;'>❌"
```

## ▼ Create Gradio Interface

```
# Dropdown for model selection
model_options = ["Simple RNN", "LSTM", "LSTM + Word2Vec"]

# Custom CSS for better styling
custom_css = """
.gradio-container {
    background-color: #f9f9f9;
    font-family: 'Arial';
}
button {
    background-color: #1e90ff;
    color: white;
    border-radius: 8px;
    padding: 10px 20px;
}
"""

# Gradio Interface
demo = gr.Interface(
    fn=predict_news,
    inputs=[
        gr.Textbox(lines=5, placeholder="Enter news article or tweet here...", label="Input Text"),
        gr.Dropdown(choices=model_options, value="LSTM + Word2Vec", label="Select Model")
    ],
    outputs=gr.HTML(label="Prediction Result"), # <-- Changed from Markdown to HTML
    title="📰 True vs Fake News Classifier",
    description="Classify whether a piece of news is true or fake using RNN, LSTM, or LSTM with Word2Vec embeddings.",
    theme="soft",
    css=custom_css,
    allow_flagging="never"
)

# Launch Gradio App
demo.launch()
```

```
→ /usr/local/lib/python3.11/dist-packages/gradio/interface.py:415: UserWarning: The `allow_flagging` parameter in `Interface` is deprecated
  warnings.warn(
It looks like you are running Gradio on a hosted Jupyter notebook. For the Gradio app to work, sharing must be enabled. Automatically
Colab notebook detected. To show errors in colab notebook, set debug=True in launch()
* Running on public URL: https://8b1163f7202e4f29cb.gradio.live

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working dir
```



**No interface is running right now**