

6CS012 – Artificial Intelligence and Machine Learning.
Lecture – 07
Representation Learning:
From Supervised {CNN} to Unsupervised {Autoencoder}.

Siman Giri {Module Leader – 6CS012}

1. What is Representation Learning?

{ Deep Learning is a Representation Learning.}

1. Introduction.

- In any Learning Task, we need to learn:
 - A good set of features to represent your data,
 - A classifier on those features,
- If you design the features manually,
 - it is called Feature Engineering, very important for Machine Learning Algorithms.
- If you do it automatically,
 - it is called Feature learning or Representation Learning, which is a big part of any Deep Learning algorithms for example Convolutional Neural Network.

1.1 How to Represent Image?

- A Classical Computer Vision Perspective:



→ Class - guarana



edge

→ Class - guarana



edge

orientation

→ Class - guarana



edge

orientation

histograms

→ Class - guarana



edge

orientation

histograms

clusters


→ Class - guarana

Challenges:

- Extract Specialized Features.
- Domain Knowledge Required.
- Easy to extract low – level representations but
- Becomes extremely difficult to design and extract high level representations.

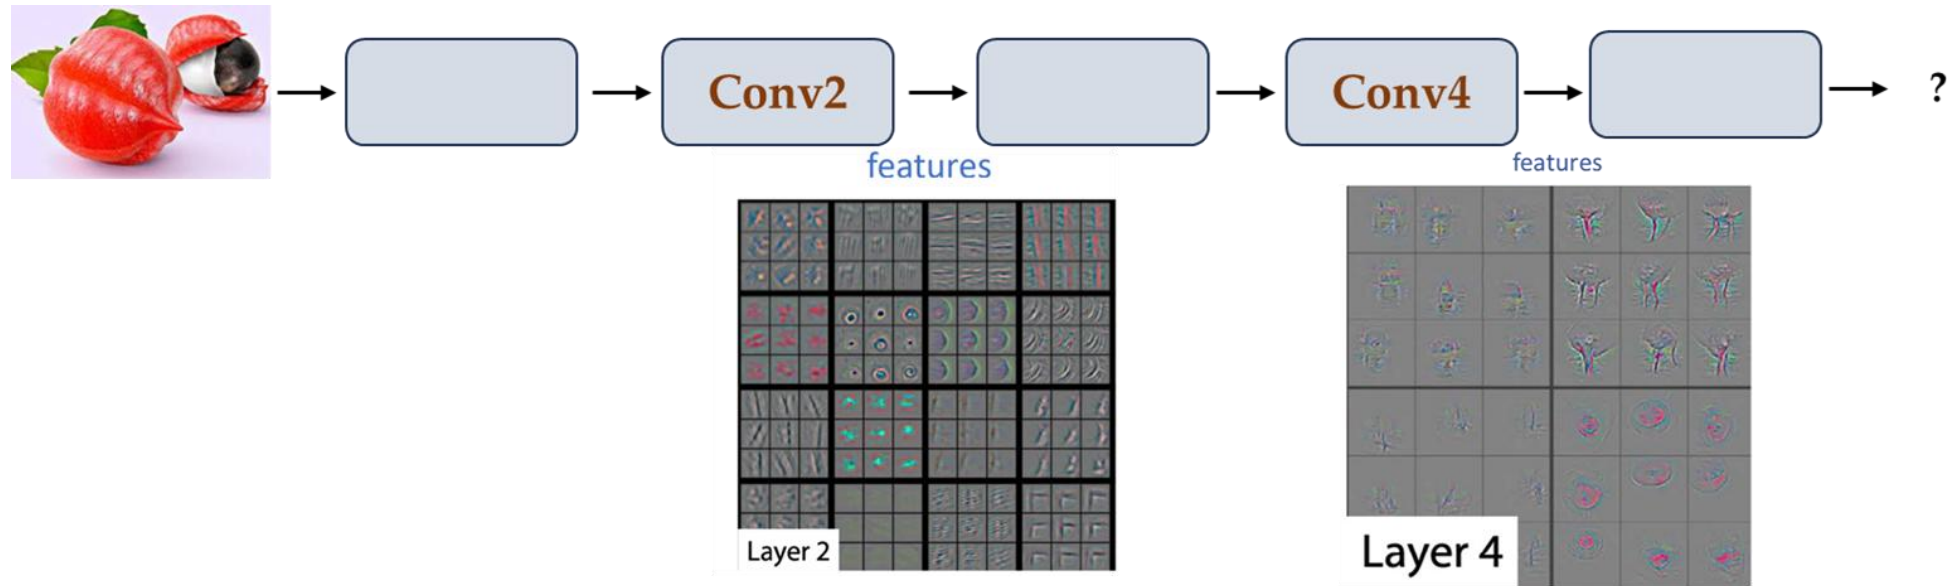
1.2 How to Represent Images with Deep Learning?

- Build a general modules instead of specialized features:
 - Compose simple modules into complex functions:



- Build multiple levels of abstractions.
- Learn from data.
- Reduce domain knowledge and feature engineering.

1.3 Multiple Levels of Representations.

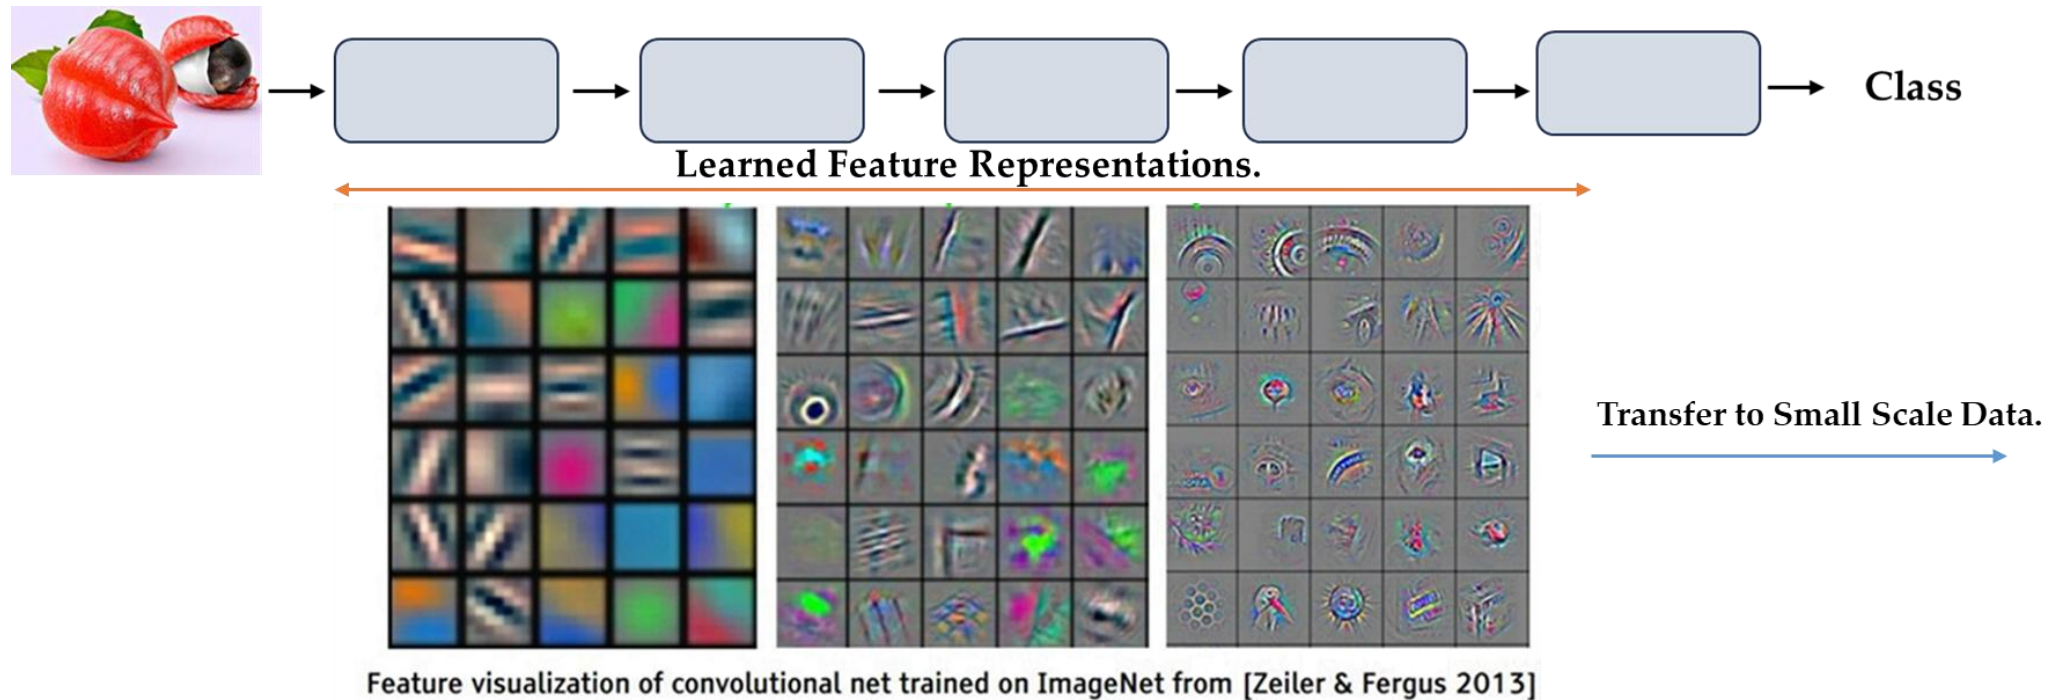


Deeper layers have “higher – level” features.

“visualizing and Understanding Convolutional Networks”, Zeiler & Fergus, ECCV 2014.

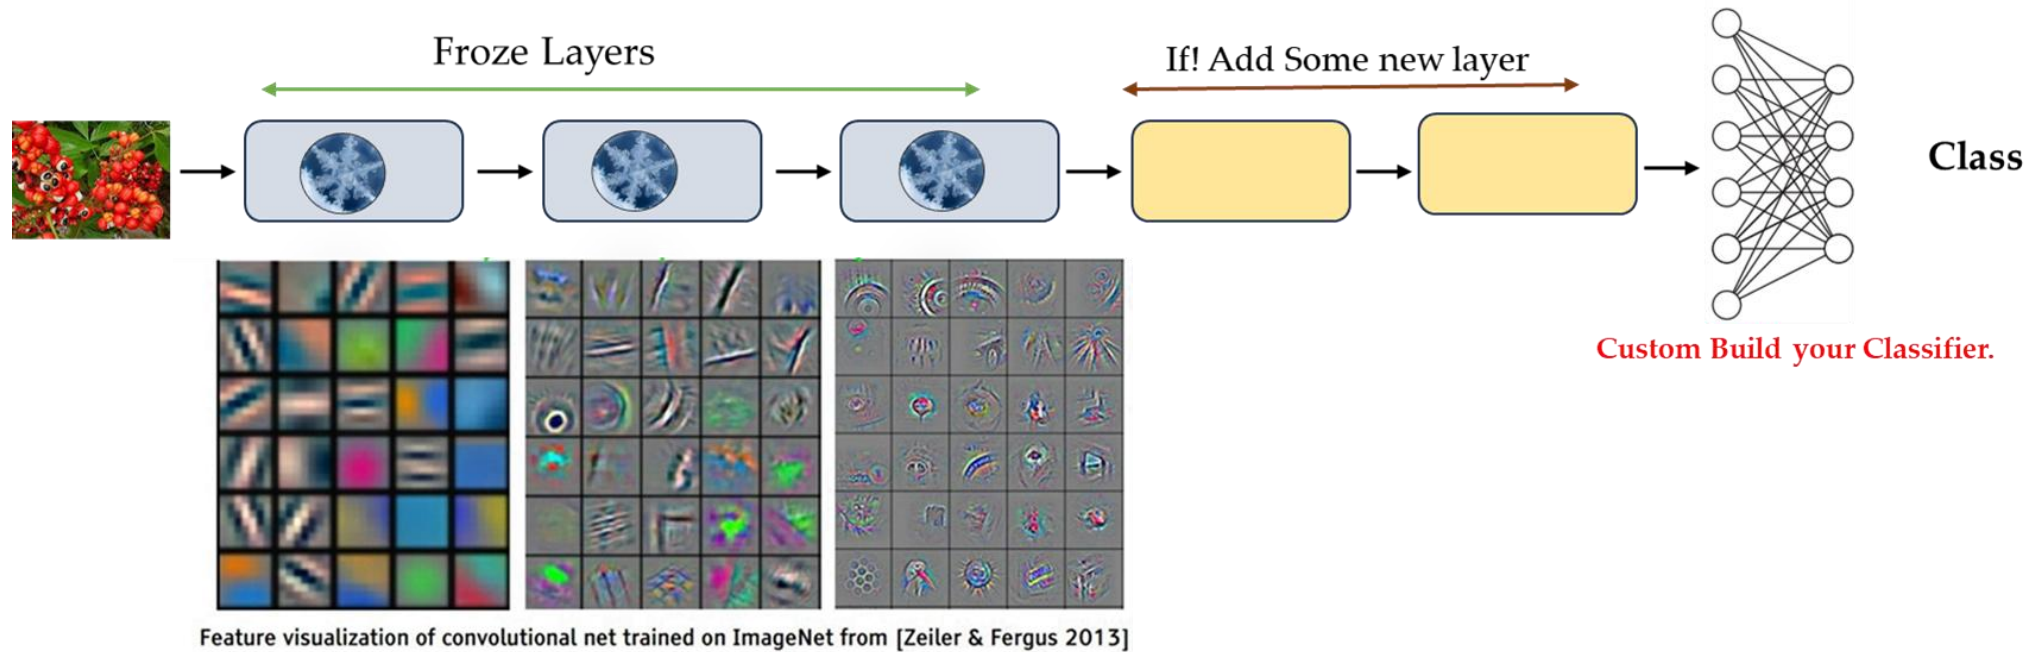
1.4 Deep Representations are Transferrable: Transfer Learning.

- Pre – train on Large Scale Model:
 - To learn general representations.
 - Train for a long time with large models on Large scale data.



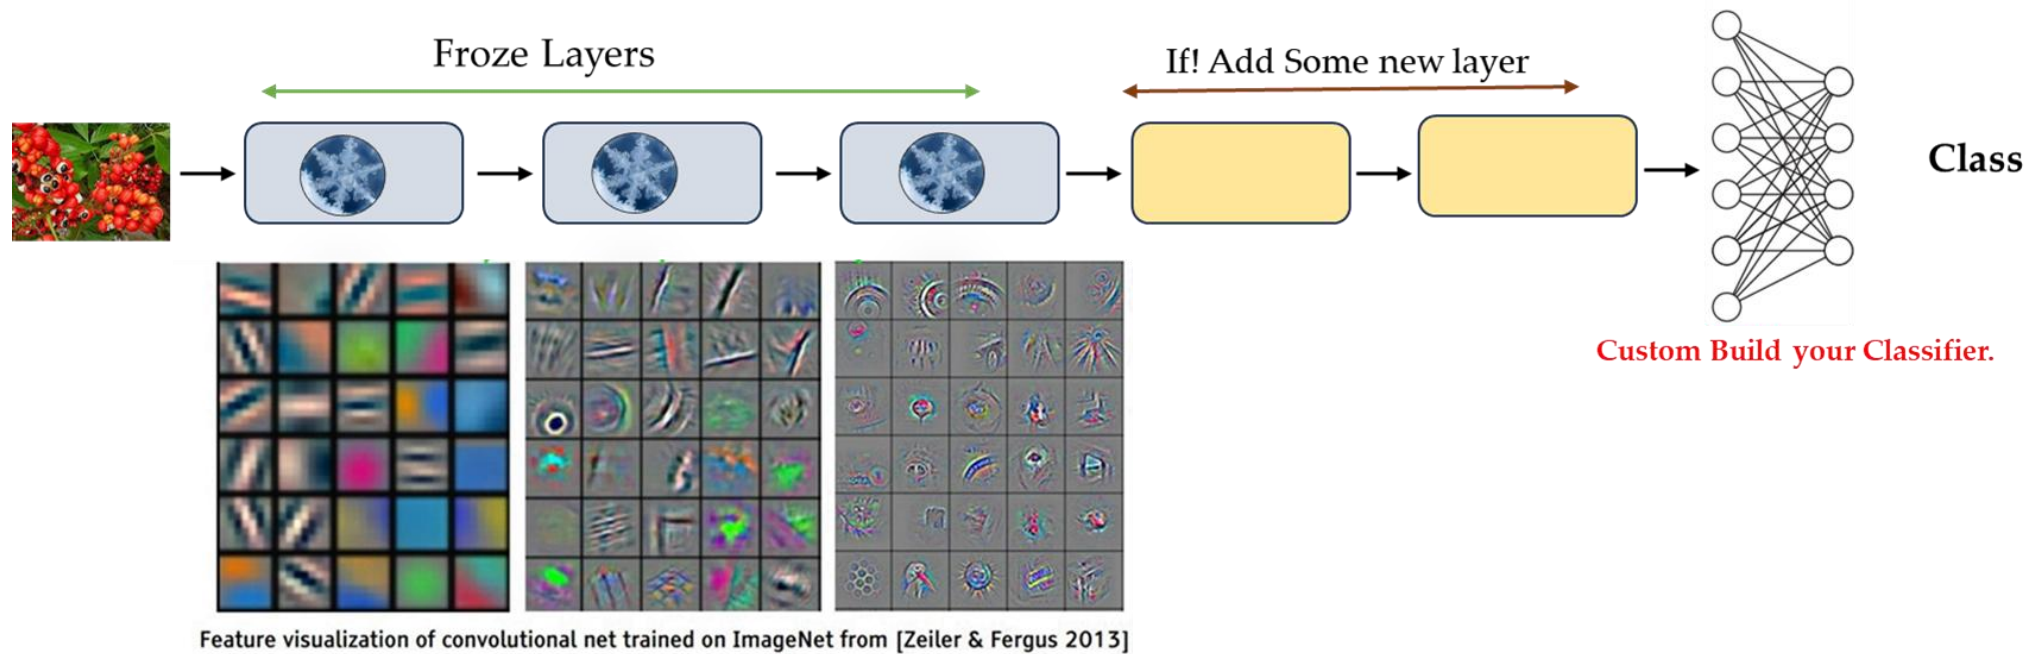
1.4.1 Deep Representations are Transferrable: Transfer Learning.

- Fine Tuning:
 - Transfer weights to specific task on small scale data
 - Train for a short time, lower learning rate.



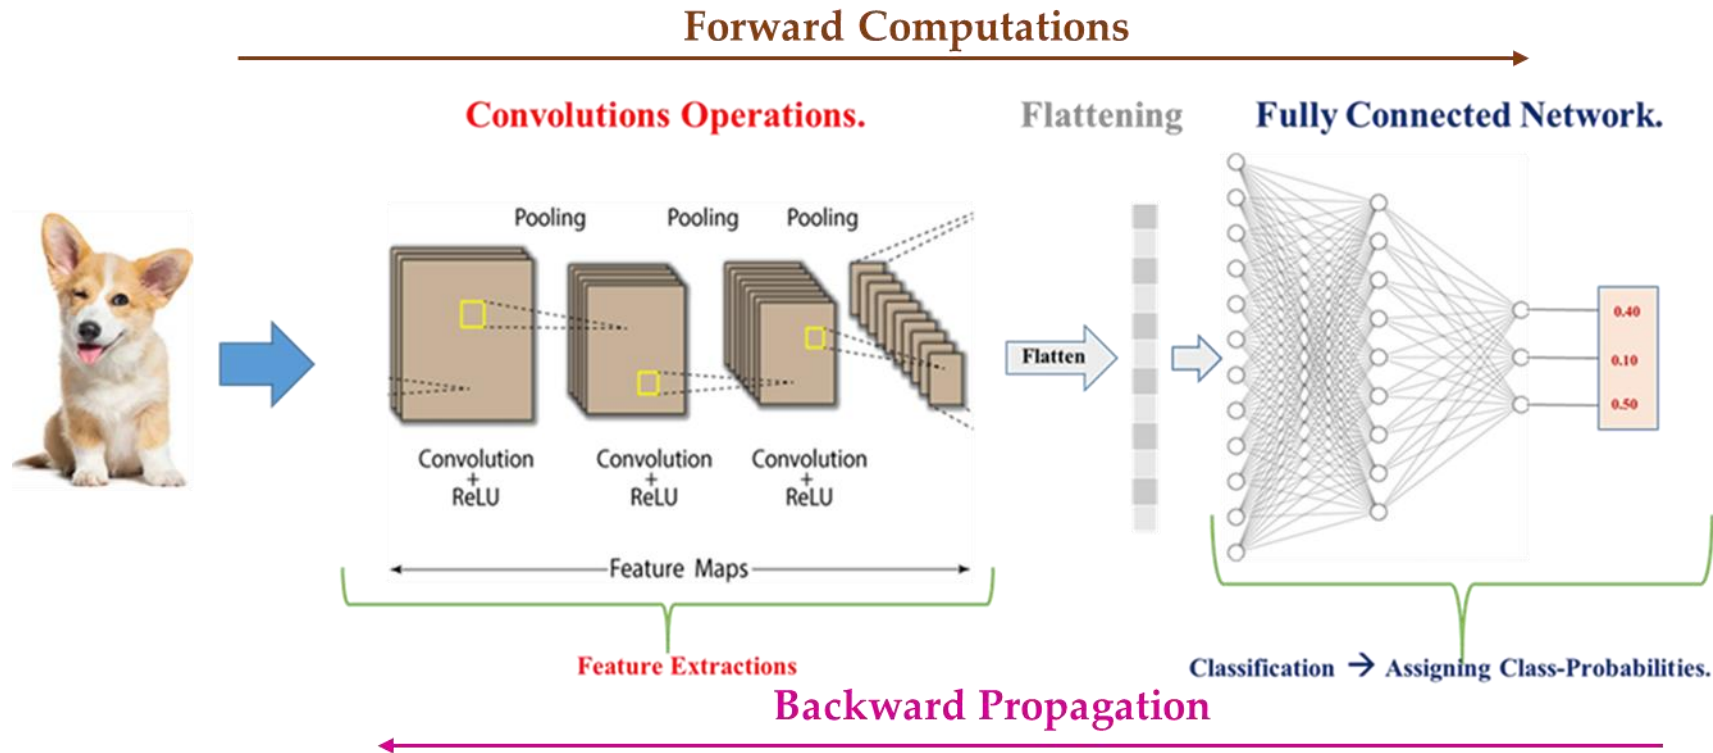
1.4.1 Deep Representations are Transferrable: Transfer Learning.

- Fine Tuning:
 - Transfer weights to specific task on small scale data
 - Train for a short time, lower learning rate.



How did we train the Model to learn those Representations?

1.5 How did we Train a CNN?



- In a typical **Convolutional Neural Network (CNN)**,
 - features are learned by **minimizing a loss function computed from the difference between the predicted and actual class labels.**
- **But can we learn meaningful features from data *without* any labels?**

1.6 Learning Objective

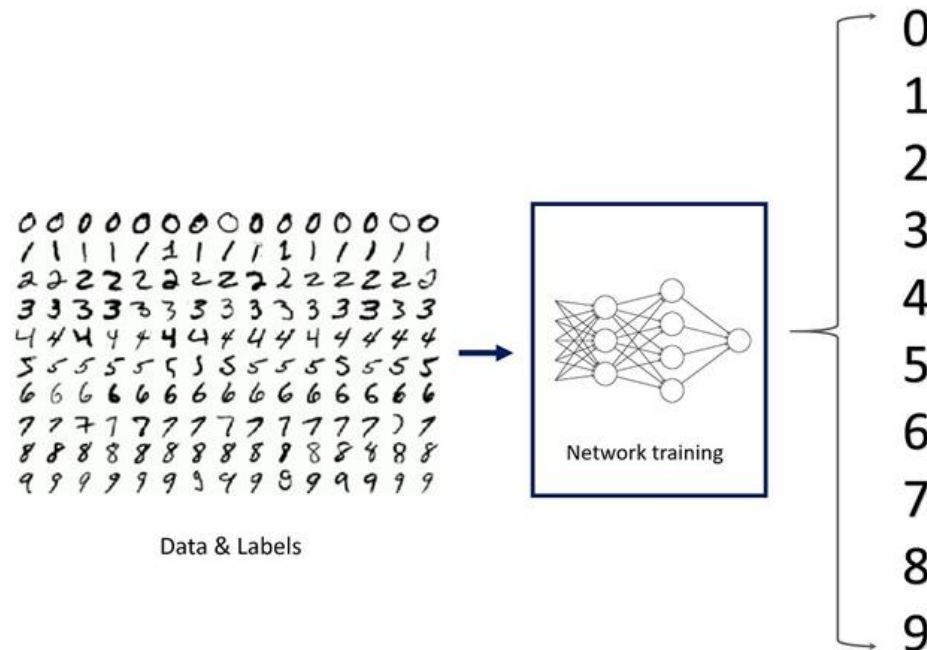
- Can we build a deep learning models
 - **which enable us to extract useful representations purely from the structure of the input data itself.**
- We will discuss
 - the underlying principles of representation learning without supervision and
 - will implement and evaluate an **autoencoder** for extracting compressed features from raw input data.

2. Background.

{Types of Learning!!!}

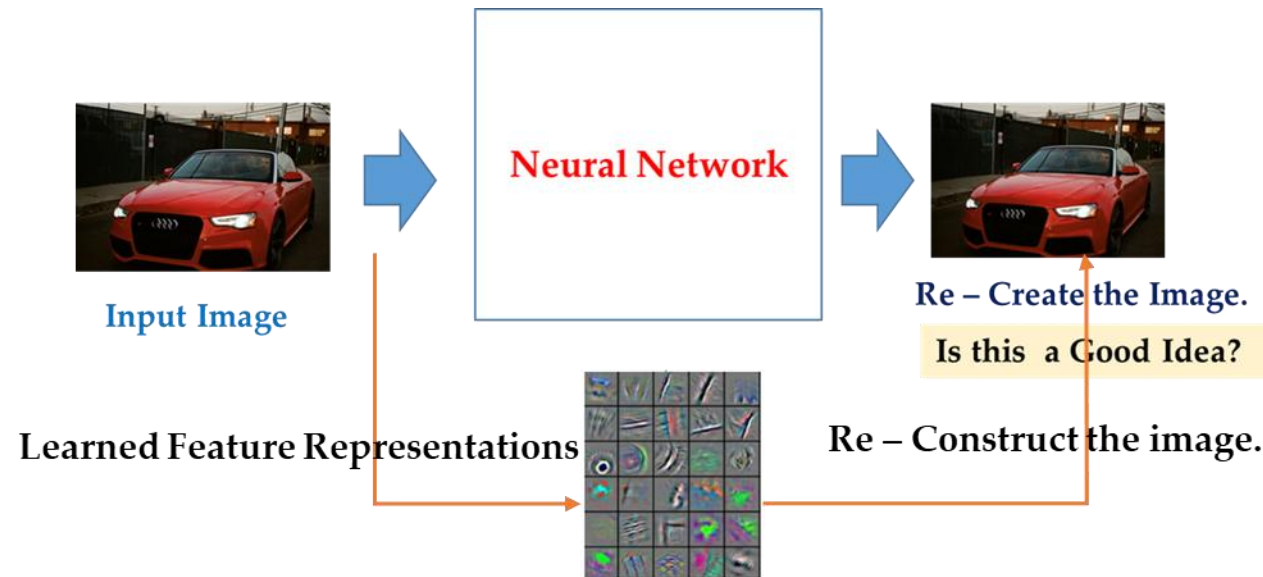
2.1 Recap: Supervised Learning.

- **Data** : $\{\mathbf{X}, \mathbf{Y}\}$ where \mathbf{X} is **data (feature)** and \mathbf{Y} is **label (target)**
- **Goal** : Learn a **Function** to map $\mathbf{X} \rightarrow \mathbf{Y}$
- **How** : Minimizing a **loss function**:
 - Loss function : divergence {difference} between \mathbf{Y} and $\hat{\mathbf{Y}}$.
- **Example**: **Classification**.



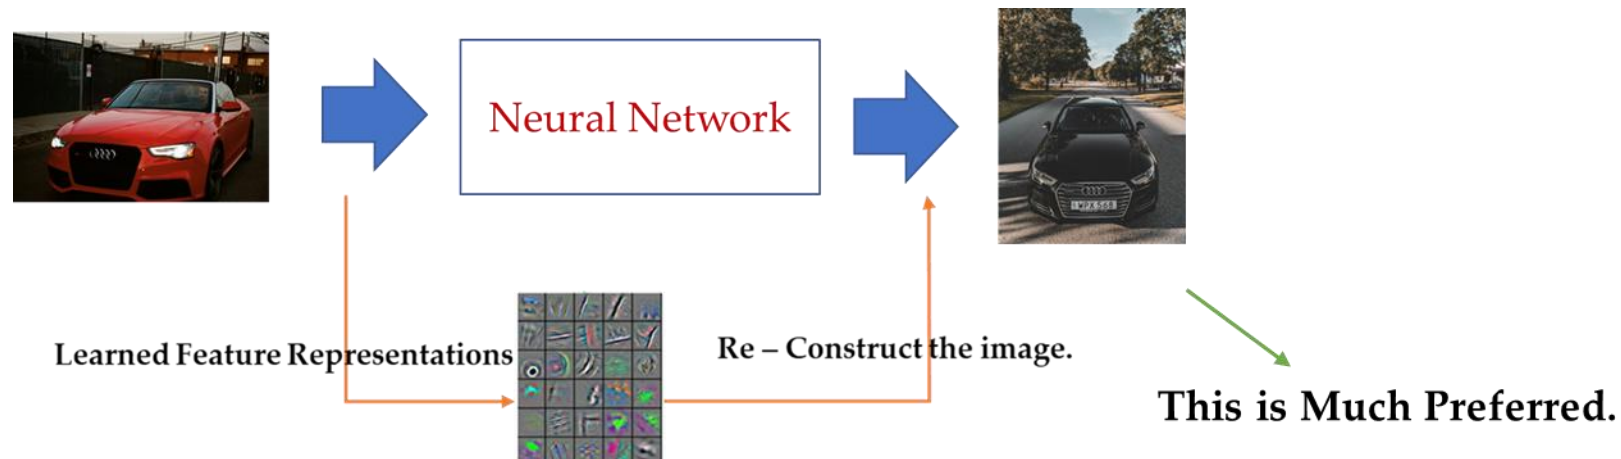
2.2 Recap: Unsupervised Learning.

- **Data** : $\{X\}$ where X is **data (feature)** and $\{Y$ is **label (target) which is not available.**
- **Goal** : Learn an underlying hidden structure {Representation} of the data. Why?
- **Example Task** : Such that we can re – create it self.
 - $X \rightarrow \tilde{X}$



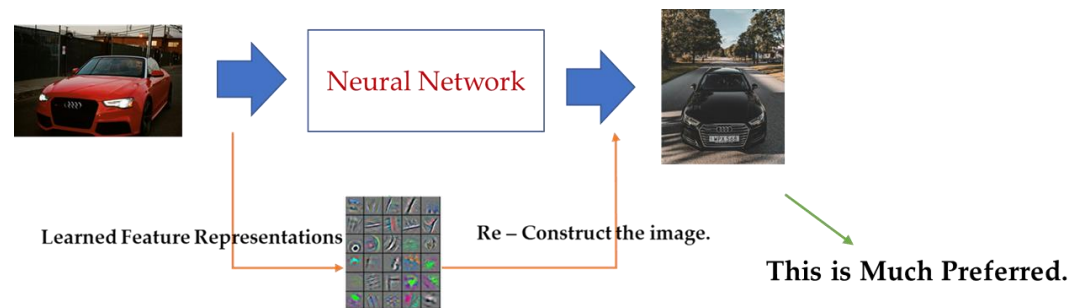
2.2.2 Unsupervised to Generative Model.

- We want to **learn an approximation of the identity function** such that
 - the reconstructed output closely resembles the original input i.e. $\hat{x} \approx x$,
- However, to ensure meaningful feature extraction rather than a trivial replication of the input,
 - the model must **learn a representation** that captures the **most salient features** of the **data**.
- This **learned representation** should **preserve essential structural and semantic information**, enabling the generation of an output that closely resembles the input while avoiding mere duplication.
 - Idea of Generative Models.



2.3 Generative Models: Idea.

- Given **training data**, generate **new samples** from **same distribution**.
 - Addresses **density estimation**, a **core problem** in unsupervised learning.
- Density Estimation:**
 - Density estimation is the problem of **reconstructing** the **probability density function** using a set of **given data points**.
 - Namely, we observe $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ as our datapoints and
 - we want to recover the underlying **probability density function** generating our **dataset**.
 - A classical approach of density estimation is the “**histogram**”.



This is what we call a Generative Models!!!!

2.3.1 Generative Models: A basic Intuition.

- Given **training data**, generate **new samples** from **same distribution**.
 - Addresses **density estimation**, a **core problem** in unsupervised learning.



Training Data $\{ \sim \mathbf{p}_{\text{data}}(\mathbf{X}) \}$



Generated Samples $\{ \sim \mathbf{p}_{\text{model}}(\mathbf{X}) \}$

- Want to learn $\mathbf{p}_{\text{model}}(\mathbf{X})$ similar to $\mathbf{p}_{\text{data}}(\mathbf{X})$ {density estimation problem.}
- How?

Disclaimer:


This should not be understood as copy pasting or duplicating the input.

2.3.2 Generative Models.

- Given **training data**, generate **new samples** from **same distribution**.
 - Addresses **density estimation**, a **core problem** in unsupervised learning.
- Want to learn $\mathbf{p}_{\text{model}}(\mathbf{X})$ similar to $\mathbf{p}_{\text{data}}(\mathbf{X})$ {density estimation problem.}
- How?



Training Data { $\sim \mathbf{p}_{\text{data}}(\mathbf{X})$ }




Generated Samples { $\sim \mathbf{p}_{\text{model}}(\mathbf{X})$ }

- Explicit density estimation:**
 - explicitly define and solve for $\mathbf{p}_{\text{model}}(\mathbf{X})$
 - Normalizing Flows**
 - Autoregressive Models – ChatGpt**
 - Diffusion Models – Dall . E**
- Implicit density estimation:**
 - learn model that can sample from $\mathbf{p}_{\text{model}}(\mathbf{X})$ without explicitly defining it.
 - Autoencoder {Topics of our Interest}**

2.4 Why Generative Models?

- Realistic samples for artwork, super-resolution, colorization, chatbot {Q n A}etc.



- Generative models of time-series data can be used for simulation and planning (reinforcement learning applications!)
- Training generative models can also**
 - enable inference of latent representations that can be useful as general features.**
 - This is what AutoEncoder does.{ This is what we will learn to do this week.}**

3. Introduction to Auto – Encoder.

3.1 Before we start:

- **Are Autoencoder a Generative Model?**

- Yes, autoencoders learn to reconstruct their input, making them a form of generative model.
- However, their generative capability is limited,
 - as they primarily reproduce the variations of the input rather than generating entirely new samples.
- Despite this limitation, autoencoders remain widely useful in applications such as
 - **dimensionality reduction**, **anomaly detection**, and **denoising**,
 - where precise feature learning and reconstruction are more important than generating novel samples.




Expectation



Reality

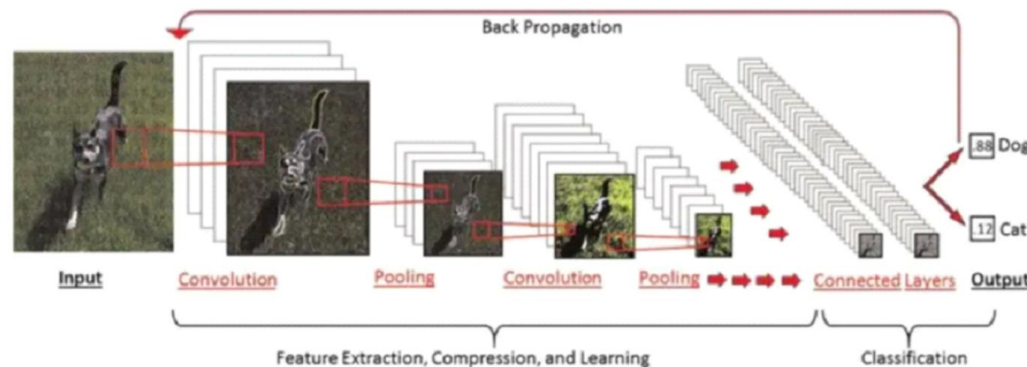
3.2 Autoencoders: Introduction.

- An autoencoder is a type of neural network architecture designed to efficiently compress (encode) input data down to its essential features,
 - then reconstruct (decode) the original input from this compressed representation.
- An autoencoder typically consists of three blocks:
 - **Encoder Layer:** to compress the input data into a compressed representations.
 - **Bottleneck layer:** or code or latent representations: to represent the compressed input.
 - **Decoder Layer:** to reconstruct the encoded image back to the original dimension.



3.3 Encoder – Decoder Architecture: Why?

- **Convolutional Neural Networks (CNNs)** traditionally follow a **canonical architecture**,
 - where the **spatial dimensions** of **feature maps** progressively decrease as **depth increases**.
 - This design is well-suited for **supervised learning**, where the primary goal is to **extract hierarchical features** and
 - **map** them to **class probabilities**.
- However, in Unsupervised learning, where **class labels are unavailable**
 - our **objective shifts from classification to learning meaningful representations**
 - that preserve essential features of the input.
 - Thus, instead of assigning class we seek to reconstruct the input itself, **necessitating a different architectural approach**.
 - This leads to Encoder – Decoder framework where:
 - **Encoder compress the input into a latent representation,**
 - **Decoder reconstructs it, enabling the model to capture the underlying structure of the data without explicit supervision.**




3.3.1 Introduction to Encoder – Decoder Architecture.

Encoder

- This process allows the encoder to capture the **most relevant information** {feature} from the inputs and produce a **fixed-size representation** of it.
- Encodes the important feature information in some latent dimension {generally smaller than input}
- Input can be image or text or others.

Decoder


- The decoder, on the other hand, takes the **encoded representation** produced by the **encoder** and **generates an output**.
- Generated output can be image or text or other.



3.3.2 Encoder – Decoder Architecture.

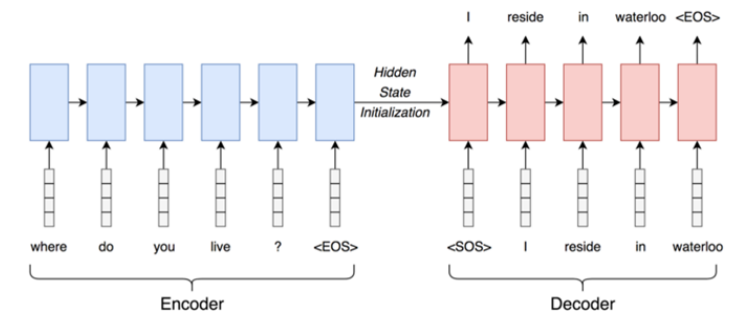
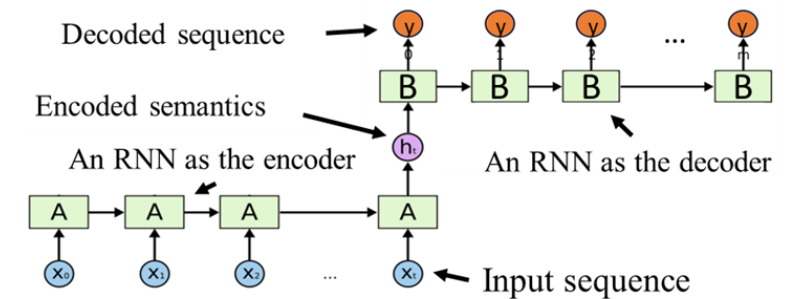
- **Training Encoder – Decoder Architecture:**

- The encoder-decoder architecture can be trained end-to-end using backpropagation and (stochastic) gradient descent to minimize a loss function that compares the predicted output to the ground truth.
- In many applications, the **{loss function}** is designed to penalize differences in the **{pixel values}** between the **predicted and ground truth images**.



3.3.3 Encoder – Decoder Architecture Application.

- Encoder – Decoder Architecture → Examples.
 - Sequence – to Sequence Model {Neural Machine Translation}
 - Encoder:
 - The encoder turns each item into a corresponding hidden vector containing the item and its context.
 - extract** and **compress** the **semantics** from the **input sequence**
 - In RNN in some variants like {LSTM or GRU is Used}
 - Decoder:
 - The decoder reverses the process, turning the vector into an output item, using the previous output as the input context.
 - generate** a **sequence** based on the **input semantics**
 - In RNN in some variants like {LSTM or GRU is Used}
 - { Week – 9's Discussion}



3.3.4 Encoder – Decoder Architecture Application.

- Encoder – Decoder Architecture → Examples.

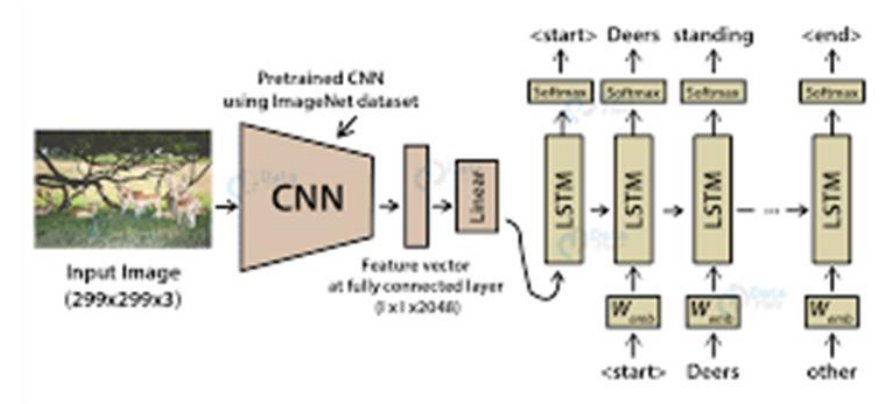
- Image Captioning

- Encoder:

- The encoder turns each image into a corresponding hidden vector containing the discriminative features.
 - extract** and **compress** the **features** from the **input images**
- An **CNN** in some variants and with transfer learning is used.


- Decoder:

- The decoder reverses the process, turning the vector into an output item, using the previous output as the input context.
 - generate** a **sequence** based on the **input features**.
- An RNN in some variants like {LSTM or GRU is Used}



3.4 Getting back to Autoencoder!!!

- Autoencoders form a very specific subset of **encoder – decoder architectures** that are trained via **unsupervised learning** to **reconstruct their own input data**.
 - **Input and output dimensions are the same.**
- Using unsupervised machine learning, autoencoders are trained to discover the **latent variables** of the input data.
- Learning $g(f(x)) = x$ (identity function) everywhere is not very useful,
 - Thus, autoencoders are designed to be unable to copy perfectly.
- Autoencoders learn useful properties of data and can prioritize which aspects of input should be copied.



Terminology Alert!!!

- **Latent Variable(Origin-Latin): “lie hidden”:**
 - In statistics:
 - latent variables are variables that **can only be inferred indirectly through a mathematical model** from **another observable variable** which can be **measured directly**.
- **Latent Representation or Space or embeddings or code:**
 - It is a vector space spanned by the latent variables.

3.5 AutoEncoder in Practice.

- Idea

- Given data \mathbf{x} (no labels) we would like to learn the functions $\mathbf{f}(\text{encoder})$ and $\mathbf{g}(\text{decoder})$ where:


- $\mathbf{f}(\mathbf{x}) = \text{act}(\mathbf{w}\mathbf{x} + \mathbf{b}) = \mathbf{z}$

- Here:

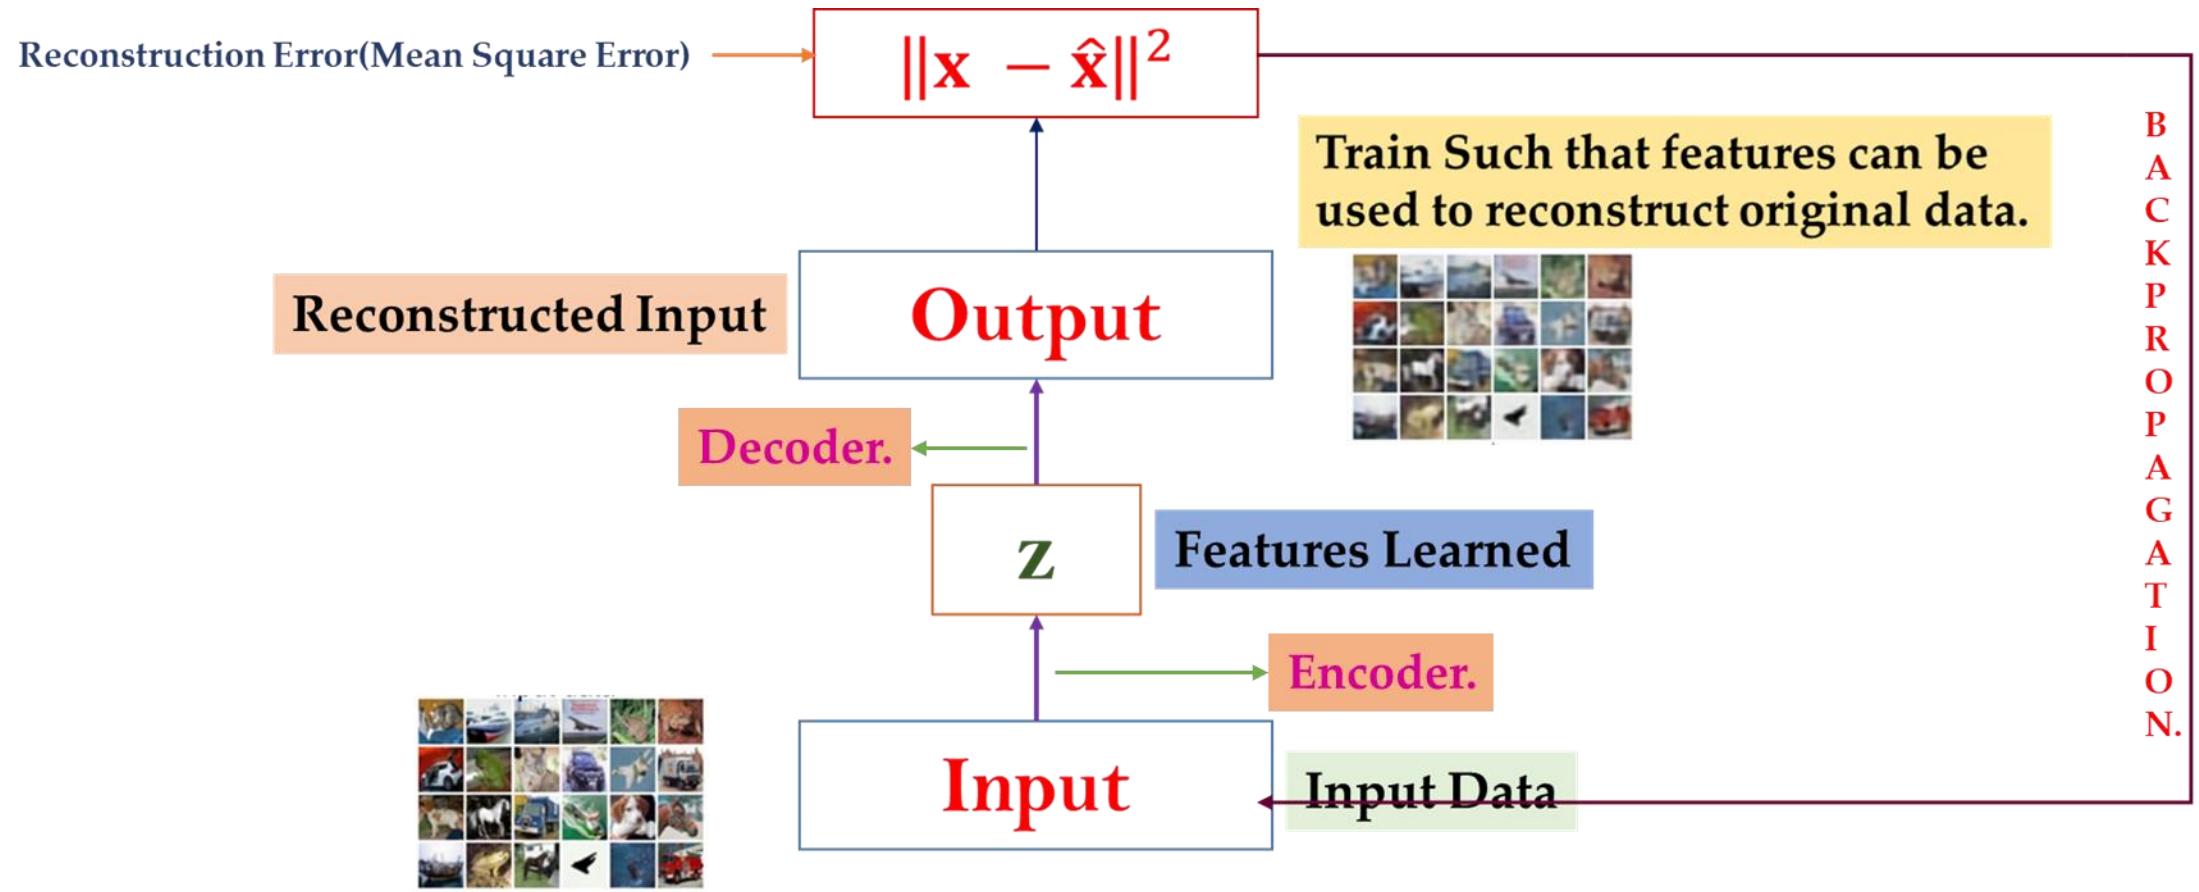
- Act – activation functions.
 - \mathbf{Z} latent representation or code.

- Here:

- $\hat{\mathbf{x}}$ is \mathbf{x} 's reconstructions.
 - $\mathbf{h}(\mathbf{x}) = \mathbf{g}(\mathbf{f}(\mathbf{x})) \rightarrow$ approximation of identity function.




3.6 Training an Autoencoder.




3.6.1 Training Auto – Encoders: Hyper – parameters.

- There are **four hyperparameters** that need to be set when training an autoencoder:
 - **Code Size (Dimension of Latent Representations):**
 - This refers to the number of nodes in the middle or bottleneck layer. By choosing a smaller size than the input dimension, we force the model to learn the most important feature representations of the data.
 - **Number of Layers:**
 - This typically refers to the number of fully connected (dense) or convolutional layers used to build the encoder and decoder.
 - The autoencoder can be made as deep as necessary, depending on the complexity of the data.



3.6.1 Training Auto – Encoders: Hyper – parameters.

- **Number of Nodes per Layer:**
 - In a standard (stacked) autoencoder, the number of nodes decreases in successive layers of the encoder and increases back in the decoder.
 - The decoder is usually **symmetric** to the encoder in terms of layer structure.
- **Loss Function:**
 - Common choices include Mean Squared Error (MSE) for continuous data and Binary Cross Entropy (BCE) for binary or normalized data.
- Autoencoders are typically trained using **backpropagation**, combined with **mini-batch gradient descent**.

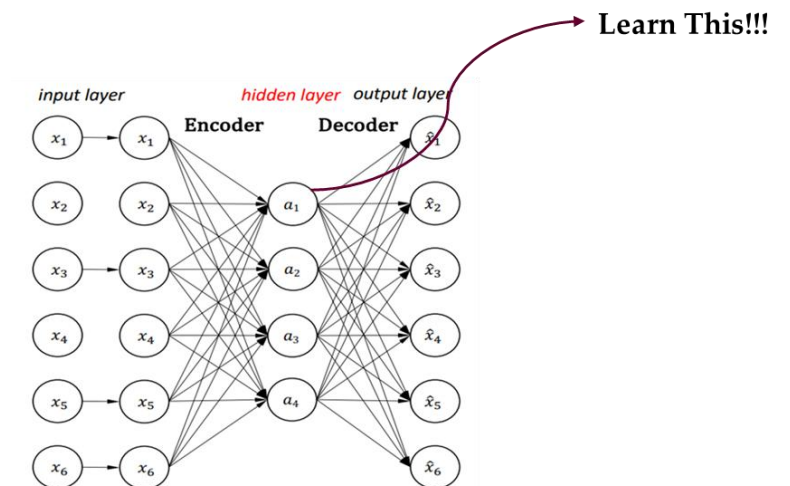


4. Designing an Auto –Encoder.

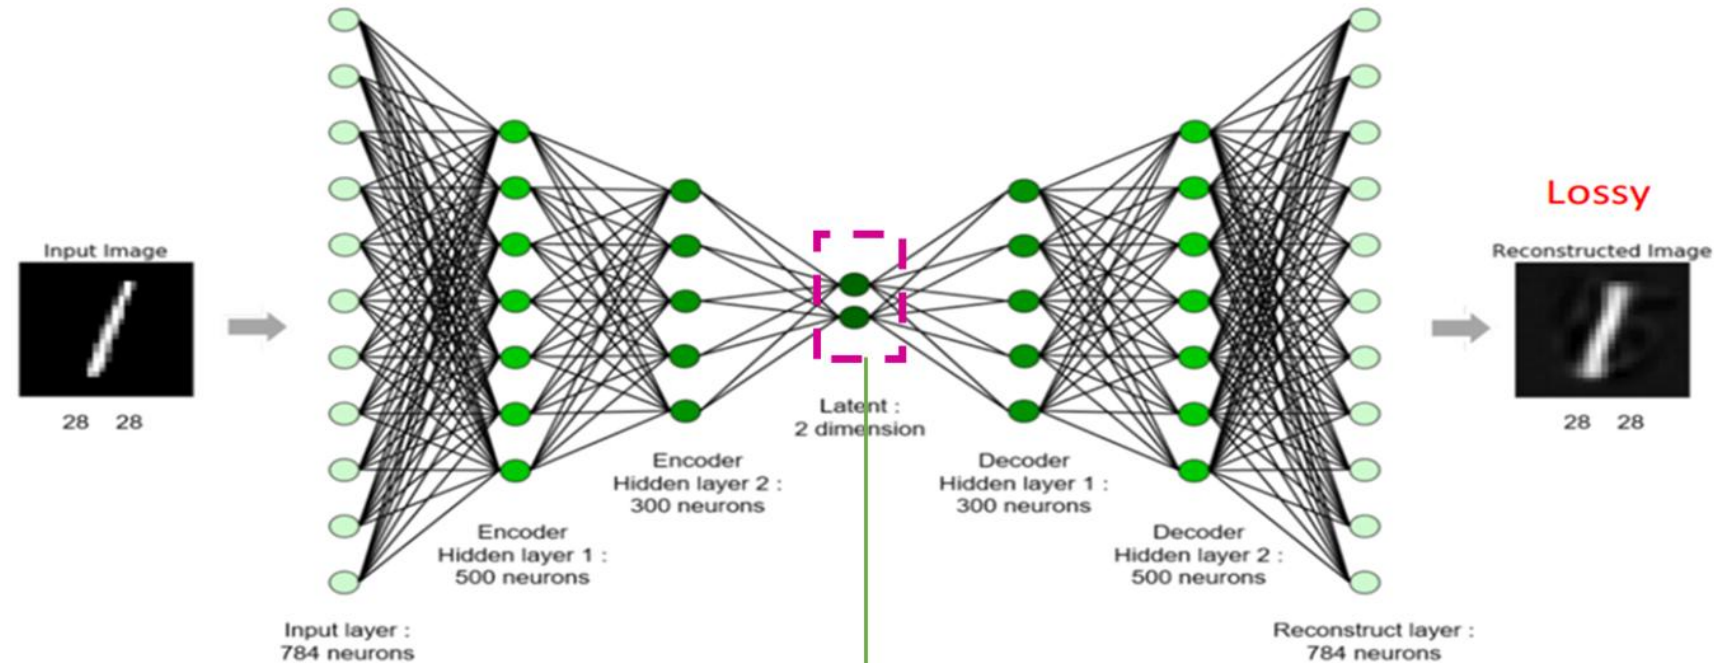
{ Putting Neural Network in Encoder and Decoder }

4.1 Vanilla Autoencoder.

- A **Vanilla Autoencoder** is the simplest form of autoencoder, composed of:
 - An **Encoder composed of Fully Connected Neural Network (Dense) layer** that maps the **input \mathbf{x}** to a **lower dimensional latent representations \mathbf{z}** .
 - A **Decoder also composed of Fully Connected Neural Network** that **reconstructs** the input from the latent code \mathbf{z} to **produce $\hat{\mathbf{x}}$** .
- **Objective:**
 - To learn a **compressed representation of the data** by mining the reconstruction loss:
 - $\mathcal{L} = \|\mathbf{x} - \hat{\mathbf{x}}\|^2$
- **Symmetric Architecture:**
 - The **decoder usually mirrors the encoder** in terms of layers and nodes.
- **Use case:**
 - Dimensionality Reduction.
 - Feature Extraction.




4.2 Vanilla Autoencoder: Dimensionality Reduction.



- Using an autoencoder with a **2-dimensional latent space**,
 - we reduce **the original 784-dimensional image representation** to just **2 dimensions**.
 - Each **image** is now encoded as a 2D point, capturing its most salient features.

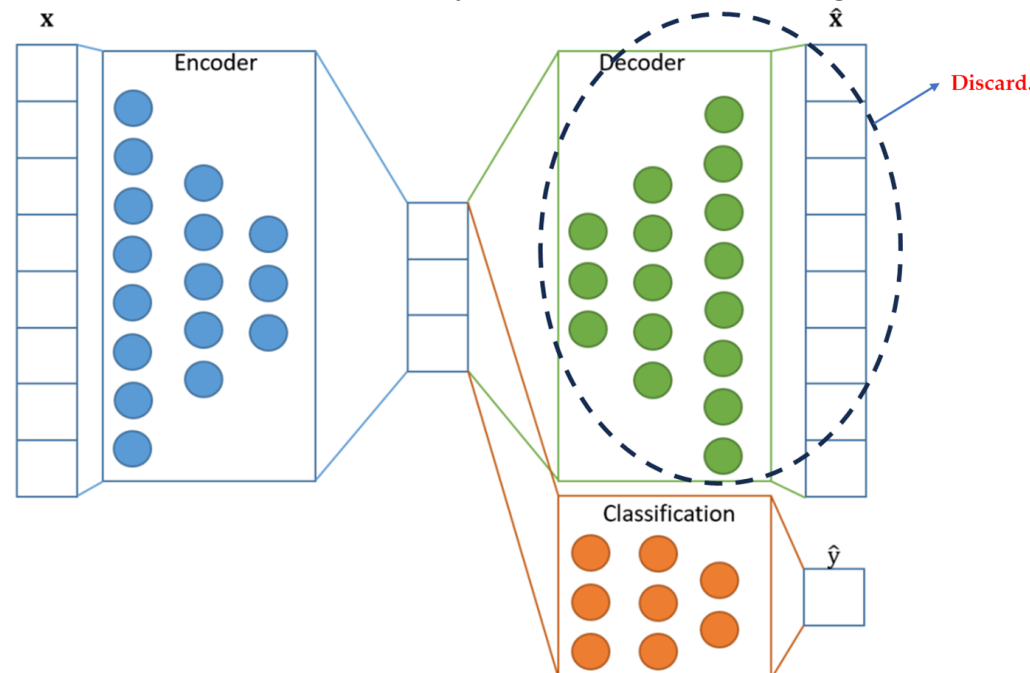
4.3 Vanilla Autoencoder: Feature Extraction.

- How it works as a Feature Extractor?
 - **Training the autoencoder:**
 - The model is trained to minimize the difference between the input and its reconstruction.
 - During this process, the **encoder** learns to capture the most informative features about the data in its latent space (i.e., the most important factors that explain the variability in the data).
 - **Using the Latent Representations:**
 - Once the model is trained, the encoder part is used to **extract features** from new data.
 - The latent space (encoded representation) becomes a **compressed version** of the input data,
 - containing the **relevant features for further tasks**.



4.3.1 Vanilla Autoencoder: Feature Extraction.

- **Feature Extraction for Downstream Tasks:**
 - **Classification:**
 - You can use the encoded features as inputs to another classifier (e.g., SVM, logistic regression).
 - **Clustering:**
 - The features can be clustered (e.g., using k-means) to identify groups within the data.
 - **Anomaly Detection:**
 - The reconstructed error can be used for anomaly detection, where a high error indicates that the input is an outlier.



4.4 Denoising Autoencoder.

- A **Denoising Autoencoder (DAE)** is a variant of the vanilla autoencoder designed to learn **robust feature representations** by reconstructing **clean inputs from noisy versions**.
- **Motivation:**
 - While vanilla autoencoders can learn compressed representations, they may **simply memorize** the input.
 - To encourage **generalization** and **robustness**,
 - denoising autoencoders are trained to **recover the original input from a corrupted version**.
- **How it works?**
 - Add **noise or corruption** to Input:
 - A stochastic corruption process (e.g. **Gaussian noise**)
 - is applied to **the input x , producing noise free \hat{x}** .
 - The original paper by Y. Bengio et. Al. (Learning Deep Architectures 2009) also suggest to use
 - **random masking similar to Dropout:**
 - During training, the **input data x** is corrupted by applying **dropout** to the input layer,
 - meaning **some parts of the input are randomly set to zero (dropped out)**.
 - This was done to **simulate the presence of missing or noisy data**.
 - {DropOut was introduced by Alex et. al in AlexNet Paper in 2012}

4.4.1 Denoising Autoencoder: Working.

- Encode the Noisy Input:
 - $\mathbf{z} = \text{Encoder}(\tilde{\mathbf{x}}) \{ \tilde{\mathbf{x}} \rightarrow \text{Noisy Input} \}$
- Decode to Reconstruct the Clean Input:
 - $\hat{\mathbf{x}} = \text{Decoder}(\mathbf{z})$
- Train to minimize Reconstruction Loss:
 - $\mathcal{L} = \|\mathbf{x} - \hat{\mathbf{x}}\|^2$
- Use Cases:
 - Image Denoising.
 - Robust Feature Extraction.

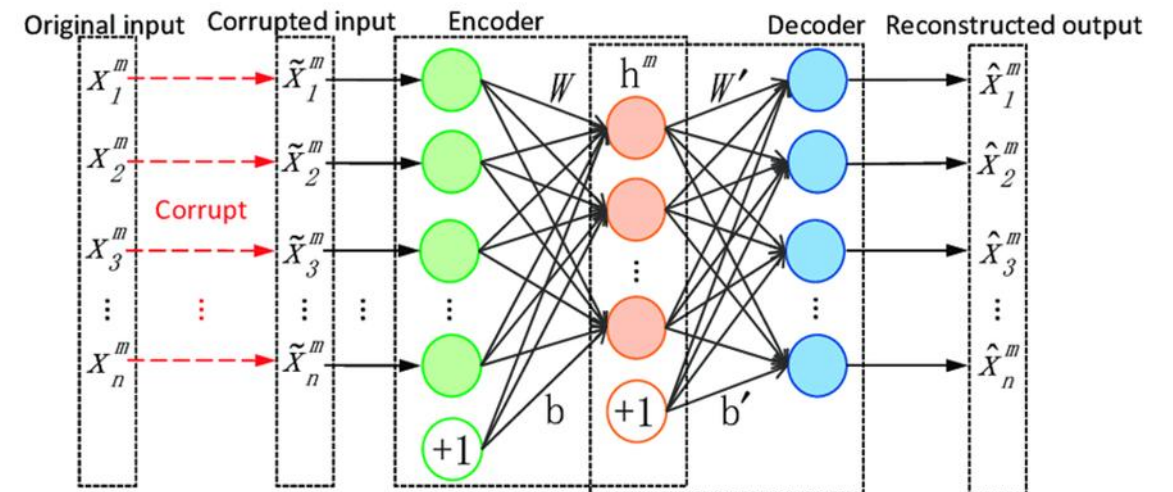


Image from Internet Subject to Copyright.

4.4.3 Denoising Autoencoder: Application.

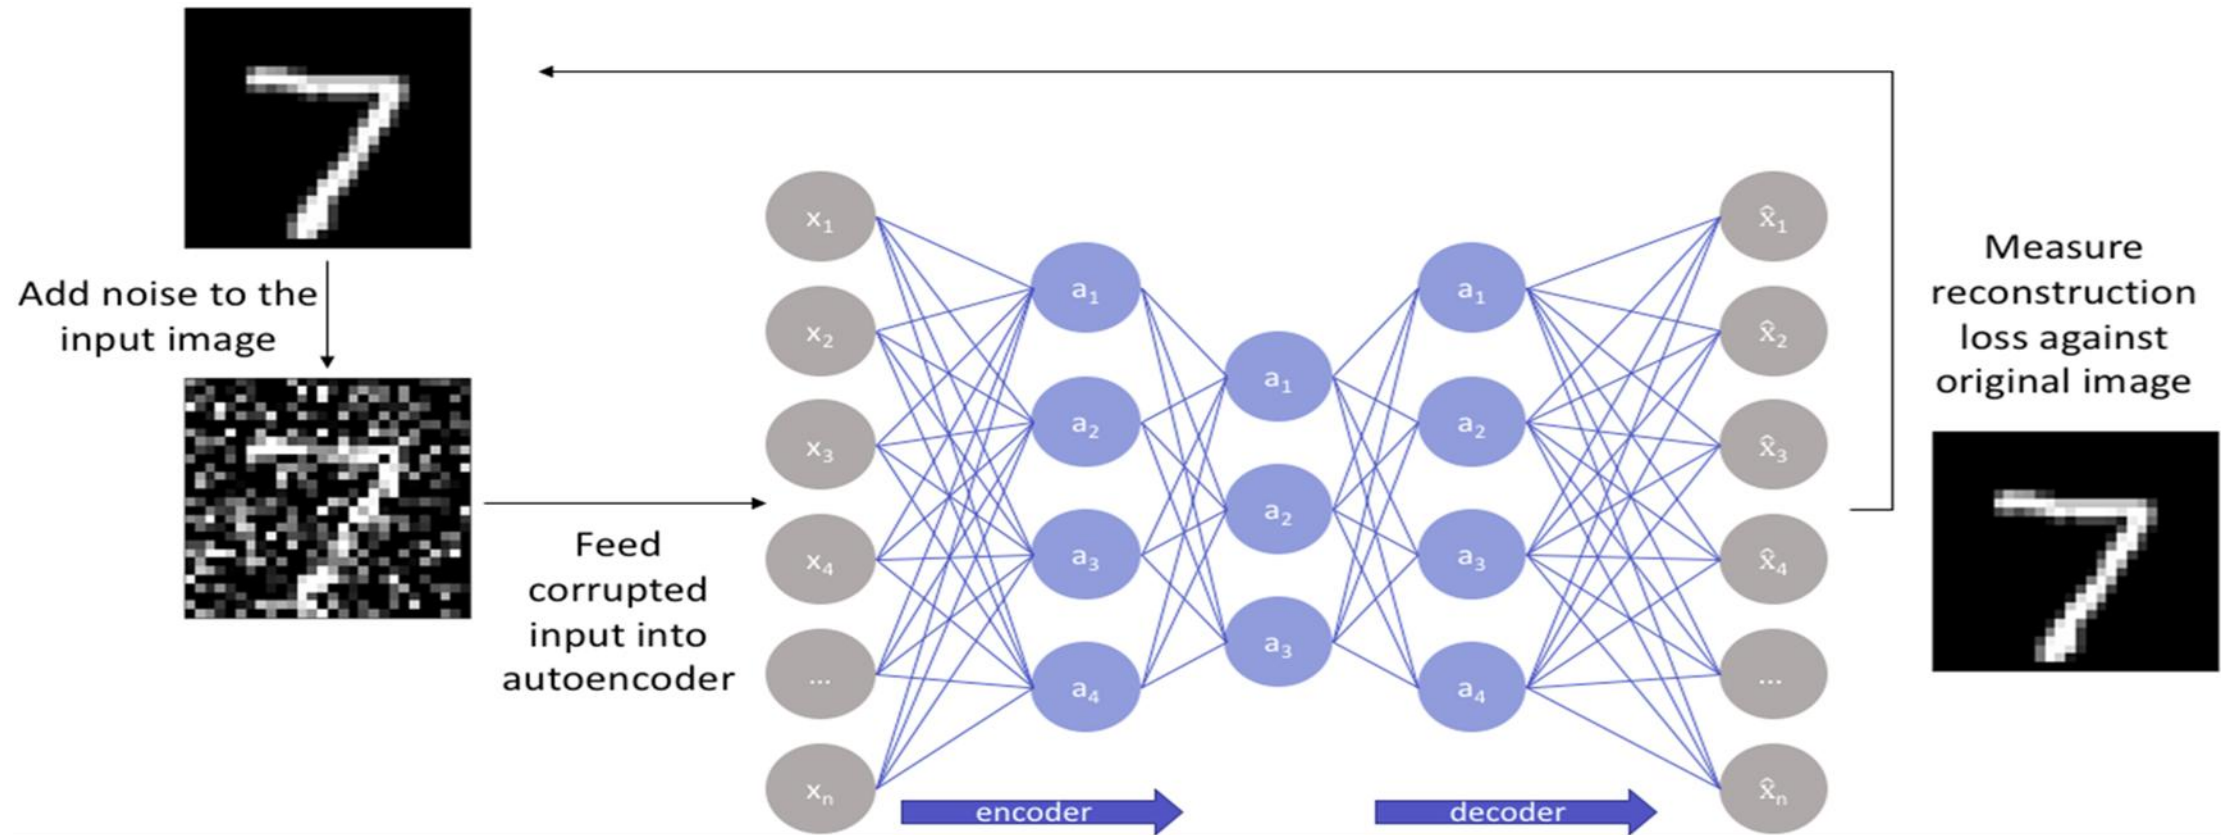




Image from Internet: Subject to copyright.

4.4.4 A Real-world application

- Water Mark Removal:



5. Convolutional Autoencoder.



5.1 Convolutional Autoencoder.

- Encoder-
 - 1 or more **convolutional** layers.
- Decoder-
 - 1 or more **Transposed convolutional** layers.

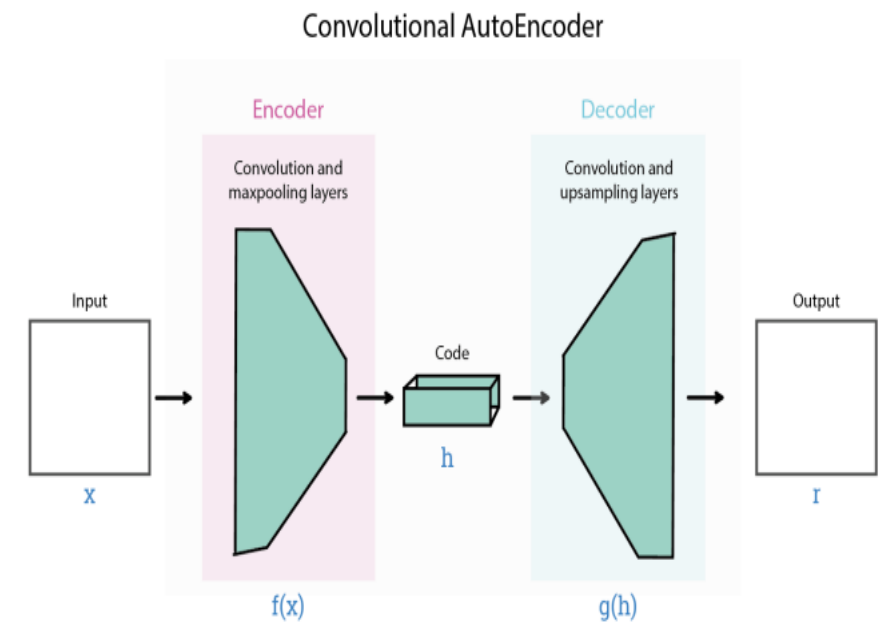



Fig. 3. General architecture of a convolutional autoencoder (CAE).

5.2 Transposed Convolutions.

- Allows us to increase the size of the **output feature map** compared to the **input feature map**.
- The output size of a transposed convolutional layer can be computed using the following formula:
 - $\text{output_size} = (\text{input_size} - 1) * \text{stride} + \text{kernel_size} - 2 * \text{padding}$
 - where:
 - **input_size**: the size of the input tensor along the spatial dimensions (width and height)
 - **stride**: the stride of the transposed convolution operation
 - **kernel_size**: the size of the transposed convolution kernel along the spatial dimensions
 - **padding**: the amount of zero padding added to the input tensor along the spatial dimensions
- Synonyms:
 - often also (**incorrectly**) called "**deconvolution**" (mathematically, deconvolution is defined as the inverse of convolution, which is different from transposed convolutions)
 - the term "unconv" is sometimes also used
 - aka "upsampling" layer.

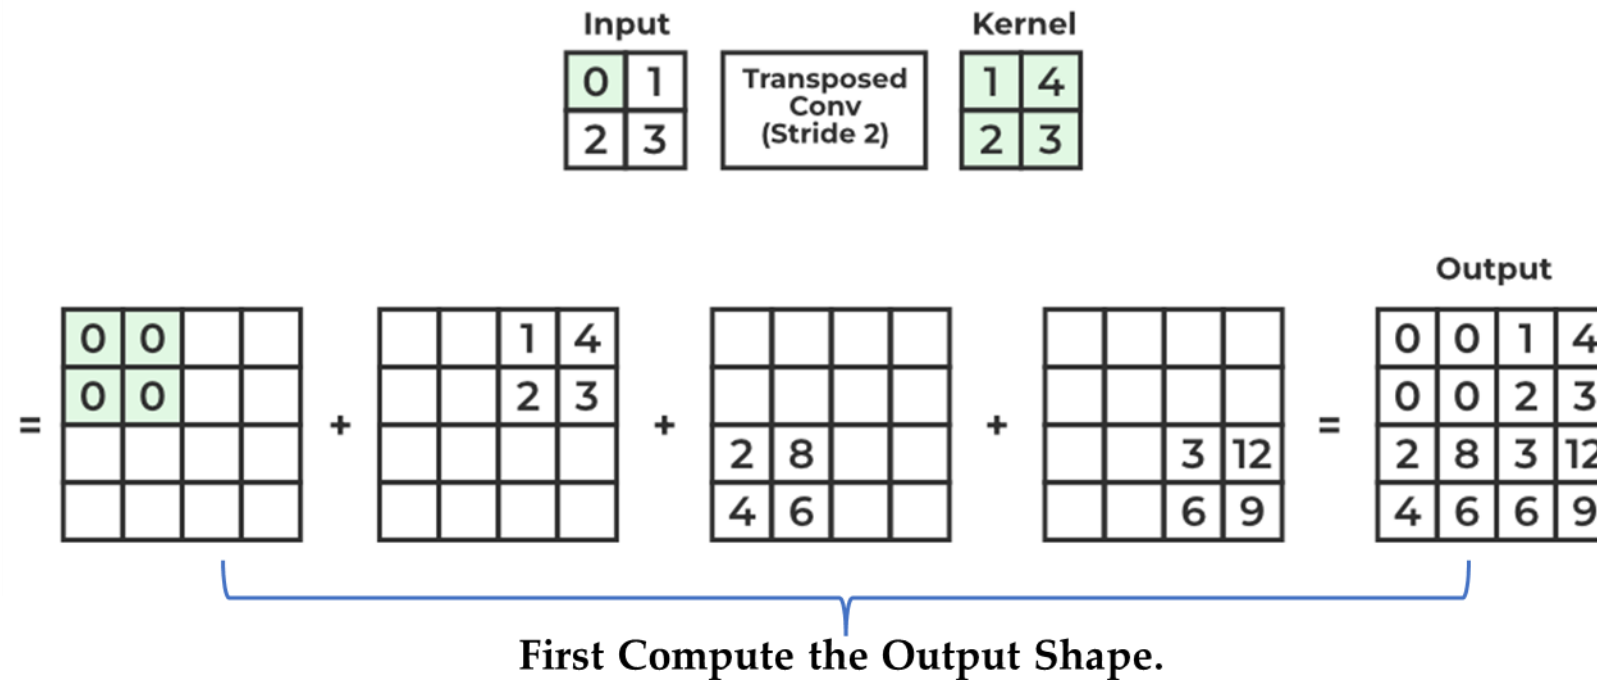
5.3 Transposed Convolutions – In Practice.

- Implementing a transposed convolutional layer can be better explained as a 4-step process:
 - Step 1:** Calculate new {hyper} parameters z and p'
 - Step 2:** Between each row and columns of the input, insert z number of zeros. This increases the size of the input to $(2*i-1) \times (2*i-1)$
 - Step 3:** Pad the modified input image with p' number of zeros
 - Step 4:** Carry out **standard convolution** on the image generated from step 3 with a stride length of 1



5.4 Transposed Convolutions-Demo.

- How Transposed Convolution Works?



5. Some Codes.

Vanilla Autoencoder.

```
# This is the dimension of the original space
input_dim = 10

# This is the dimension of the latent space (encoding space)
latent_dim = 2

encoder = Sequential([
    Dense(128, activation='relu', input_shape=(input_dim,)),
    Dense(64, activation='relu'),
    Dense(32, activation='relu'),
    Dense(latent_dim, activation='relu')
])

decoder = Sequential([
    Dense(64, activation='relu', input_shape=(latent_dim,)),
    Dense(128, activation='relu'),
    Dense(256, activation='relu'),
    Dense(input_dim, activation=None)
])
```

```
autoencoder = Model(inputs=encoder.input, outputs=decoder(encoder.output))
autoencoder.compile(loss='mse', optimizer='adam')
```

CAE.

```
input = layers.Input(shape=(28, 28, 1))

# Encoder
x = layers.Conv2D(32, (3, 3), activation="relu", padding="same")(input)
x = layers.MaxPooling2D((2, 2), padding="same")(x)
x = layers.Conv2D(32, (3, 3), activation="relu", padding="same")(x)
x = layers.MaxPooling2D((2, 2), padding="same")(x)

# Decoder
x = layers.Conv2DTranspose(32, (3, 3), strides=2, activation="relu", padding="same")(x)
x = layers.Conv2DTranspose(32, (3, 3), strides=2, activation="relu", padding="same")(x)
x = layers.Conv2D(1, (3, 3), activation="sigmoid", padding="same")(x)

# Autoencoder
autoencoder = Model(input, x)
autoencoder.compile(optimizer="adam", loss="binary_crossentropy")
autoencoder.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 28, 28, 1)]	0

conv2d (Conv2D)	(None, 28, 28, 32)	320

max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0

conv2d_1 (Conv2D)	(None, 14, 14, 32)	9248

max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 32)	0

conv2d_transpose (Conv2DTran)	(None, 14, 14, 32)	9248

conv2d_transpose_1 (Conv2DTr)	(None, 28, 28, 32)	9248

conv2d_2 (Conv2D)	(None, 28, 28, 1)	289
=====		
Total params: 28,353		
Trainable params: 28,353		
Non-trainable params: 0		

At the end!!!!

- Autoencoders learn data representation in an **unsupervised**/ **self-supervised** way.
 - Learned features are able to capture salient properties of data
- Different with vanilla autoencoder, in sparse autoencoder, the number of hidden units can be greater than the number of input variables.
 - **Under - complete** and **Over - complete** Architectures.
- Can also be stacked to create → deep autoencoders.
- **“You will implement Autoencoders in Tutorial → Please come with laptop.”**

Thank You